


Graphical Abstract

Random Matrix Filtering and Planar Financial Networks: A Retrospective Assessment for Volatility Forecasting in Complex Financial Markets

Andrey V. S. Souza, Bruno L. Dalmazo, Giancarlo Lucca, Eduardo N. Borges, Richard F. Pinto, Fabian C. Cardoso, Viviane L. D. de Mattos, Rafael A. Berri



Spectral filtering exposes collective and localized dependency structure; the final model comparison is retrospective because structural descriptors use the complete synchronized sample.

Highlights

Random Matrix Filtering and Planar Financial Networks: A Retrospective Assessment for Volatility Forecasting in Complex Financial Markets

Andrey V. S. Souza, Bruno L. Dalmazo, Giancarlo Lucca, Eduardo N. Borges, Richard F. Pinto, Fabian C. Cardoso, Viviane L. D. de Mattos, Rafael A. Berri

- Random-matrix spectra identify market, group and noise-compatible modes.
- Planar filtering exposes localized sectoral dependency patterns.
- Within-sector correlations exceed between-sector correlations on average.
- Full-sample network descriptors yield modest ex-post model differences.
- Compact neural models extend the machine-learning comparison.

Random Matrix Filtering and Planar Financial Networks: A Retrospective Assessment for Volatility Forecasting in Complex Financial Markets

Andrey V. S. Souza^{a,*}, Bruno L. Dalmazo^a, Giancarlo Lucca^a, Eduardo N. Borges^a, Richard F. Pinto^a, Fabian C. Cardoso^b, Viviane L. D. de Mattos^a and Rafael A. Berri^a

^aFederal University of Rio Grande (FURG), Rio Grande, Rio Grande do Sul, Brazil

^bUniversity of Rio Verde (UniRV), Rio Verde, Goiás, Brazil

ARTICLE INFO

Keywords:

Econophysics
Random matrix theory
Financial networks
Planar maximally filtered graph
Volatility forecasting
Machine learning
Brazilian stock market

ABSTRACT

We study the dependency structure of Brazilian equities traded on B3 and assess, retrospectively, whether structural descriptors are associated with realized-volatility model performance. Using adjusted daily prices, we construct returns for a 58-equity historical universe and estimate random matrix theory (RMT), clustering and network quantities from a complete synchronized panel spanning 23 August 2006 to 15 August 2024. The Marčenko–Pastur benchmark separates a dominant market mode, candidate group-level modes and a residual noise-compatible component; five eigenvalues exceed the upper bound, and the largest accounts for about 37% of the correlation-matrix trace. Minimum Spanning Tree (MST) and Planar Maximally Filtered Graph (PMFG) representations show that market-mode removal lowers average edge correlation while preserving localized patterns; the group-mode PMFG has lower mean edge correlation but higher clustering than the original PMFG. We then conduct a chronological train-validation-test assessment with machine-learning and simple neural models. RMT and network descriptors are estimated from the full synchronized sample, so the experiment identifies an ex-post structural association rather than a real-time forecasting gain. The contribution is an integrated, reproducible framework for interpreting collective B3 market structure and specifying a future point-in-time volatility experiment.

1. Introduction

Financial markets are often studied as complex systems in which heterogeneous agents, institutional constraints, liquidity cycles, information flows and feedback mechanisms interact across multiple time scales [1, 2]. Even when the analysis is restricted to daily asset returns, market variation also has a relational component: how assets move together across the cross-section, summarized by a correlation matrix. Empirical correlations can contain broad co-movement, localized organization and sampling variation; Random Matrix Theory (RMT) provides a benchmark for separating components compatible with random estimation effects from components that merit structural interpretation [3, 4, 5].

The distinction matters because financial risk is inherently multivariate. The volatility of an individual asset matters, but portfolio risk, contagion, diversification and systemic vulnerability depend on how assets co-move [6]. During calm periods, correlations may appear moderate and dispersed. During stress periods, correlations often increase, reducing diversification benefits and lowering the effective dimensionality of the market [7, 8].

Brazil is a useful setting because B3 (Brasil, Bolsa, Balcão), the country's exchange and financial-market infrastructure provider¹, contains firms from commodity-linked, financial, utility, industrial, consumer, real-estate and telecommunications activities. This diversity makes it possible to ask whether a broad market component coexists with localized dependency patterns. Brazilian studies have examined interbank networks [9], multiscale correlations [10], agrarian cross-correlations [11] and equity-market topology [12]. The specific gap addressed here is narrower: to our

*Corresponding author.

✉ andreyvinicius@furg.br (A.V.S. Souza); dalmazo@furg.br (B.L. Dalmazo); giancarlo.lucca@furg.br (G. Lucca); eduardoborges@furg.br (E.N. Borges); richardpinto@furg.br (R.F. Pinto); fabian@unirv.edu.br (F.C. Cardoso); vivianemattos@furg.br (V.L.D.d. Mattos); rafaelberri@furg.br (R.A. Berri)

ORCID(s): 0009-0006-4177-758X (A.V.S. Souza); 0000-0002-6996-7602 (B.L. Dalmazo); 0000-0002-3776-0260 (G. Lucca); 0000-0003-1595-7676 (E.N. Borges); 0009-0007-0176-3383 (R.F. Pinto); 0000-0002-2842-0387 (F.C. Cardoso); 0000-0002-3512-6290 (V.L.D.d. Mattos); 0000-0002-3812-4186 (R.A. Berri)

¹<https://www.b3.com.br/>

knowledge, few B3 equity studies jointly combine RMT filtering, planar financial networks and a clearly qualified set of retrospective model assessments.

Correlation matrices are noisy, especially when the number of assets is not negligible relative to the number of observations. Random Matrix Theory (RMT) offers a benchmark for this problem by comparing the eigenspectrum with the spectrum expected from random correlation matrices [13, 3, 4]. Eigenvalues lying inside the Marčenko–Pastur band are compatible with random sampling noise under the benchmark model, whereas eigenvalues outside the band indicate candidate collective modes that deserve further interpretation. Subsequent work has extended RMT-based methods for cleaning large covariance matrices [5]. We use RMT as an intermediate layer for reconstructing broad market, group-level and noise-compatible components before graph construction, not as a mechanical classifier of economic sectors.

Network filtering gives an additional representation of the same dependency system. By transforming correlations into distances, financial networks retain selected dependency links among assets and make it possible to inspect market backbones, local clusters and central nodes. The Minimum Spanning Tree (MST) emphasizes the sparsest connected structure [14], while the Planar Maximally Filtered Graph (PMFG) retains a denser planar representation with richer local connectivity [15, 16]. We use these constructions to ask whether strong dependencies are organized by sectors, whether central nodes change after market-mode removal and whether the intermediate-scale organization of the market contains information beyond average correlation. Filtered networks are also relevant for risk analysis, where peripheral assets can offer better diversification properties than central hubs [17, 18].

We develop an integrated econophysics workflow for Brazilian equities traded on B3. It documents stylized facts and correlation structure; applies RMT, clustering and filtered networks to interpret the dependency matrix; and then conducts chronological retrospective model assessments. The structural descriptors in those experiments are estimated from the full synchronized sample. Consequently, the experiments ask whether those descriptors are associated with model performance, not whether they deliver information available at a historical forecast origin.

The methodological novelty is the controlled connection of these layers into one auditable experiment. We hold the asset universe, ranking rule and network algorithms fixed while comparing the original correlation structure with the structure that remains after the broad market mode is removed. This design treats the leading eigenmode as a collective synchronization channel and the residual topology as a mesoscopic organization of the market. The resulting comparison is intended to reveal qualitative changes in clustering, hubs and sectoral connectivity, rather than to propose a new RMT or graph algorithm. The implementation, configuration files and reproducibility documentation are available in the accompanying research-software repository [19].

The paper makes four contributions:

1. We build a reproducible econophysics workflow for B3 equities using adjusted daily prices with raw historical coverage from 1998 to 2025, including return construction, stylized-fact checks, correlation analysis, sectoral comparison and rolling synchronization.
2. We apply Random Matrix Theory to the complete 58-asset synchronized panel (23 August 2006–15 August 2024) and reconstruct market, group/sector and noise-compatible components of the empirical correlation matrix, yielding a spectral interpretation of Brazilian equity dependencies.
3. We compare MST and PMFG networks built from original and RMT-filtered matrices, with emphasis on hub reorganization, clique structure, clustering, sectoral retention and subsector-level aggregation.
4. We conduct a chronological, nested-feature-set retrospective assessment of whether full-sample RMT and network descriptors are associated with realized-volatility model performance. This motivates, but does not replace, a point-in-time forecasting design with recursively estimated features.

The article is organized as follows. Section 2 establishes the theoretical background and conceptual foundations. Section 3 reviews the relevant literature. Section 4 presents the proposed reproducible integrated framework, including the data and asset universe, spectral decomposition, filtered networks and retrospective model assessments. The subsequent sections report the empirical results for stylized facts, RMT, clustering, network topology, subsector aggregation and volatility models, followed by discussion, limitations and conclusion.

2. Theoretical Background and Conceptual Framework

This section establishes the theoretical, mathematical and econophysics foundations of the research pipeline. We formalize cross-sectional financial dependencies and finite-sample estimation noise, Random Matrix Theory (RMT)

spectral filtering, topological planar graphs (MST and PMFG), and the principles of realized volatility modeling in complex networks.

2.1. Cross-Sectional Dependencies and Complex Market Dynamics

Financial markets operate as complex interconnected systems where asset prices continuously interact through macroeconomic news, sector trends, liquidity cycles, and investor behavior [1, 20]. Rather than analyzing individual stocks in isolation, modern risk management focuses on understanding how assets move together across the entire market cross-section.

To compare price changes consistently across different companies, prices are converted into daily logarithmic returns:

$$r_{i,t} = \ln \left(\frac{P_{i,t}^{\text{adj}}}{P_{i,t-1}^{\text{adj}}} \right), \quad (1)$$

where $P_{i,t}^{\text{adj}}$ is the dividend- and split-adjusted closing price of asset $i \in \{1, \dots, N\}$ on trading day $t \in \{1, \dots, T\}$ [21, 22]. Across global stock markets, return series share several well-established universal behaviors, known as *stylized facts* [2]:

- **Heavy Tails (Non-Gaussianity):** Extreme positive or negative price swings occur much more frequently than predicted by a standard normal curve (leptokurtosis) [23, 24].
- **Volatility Clustering:** While raw daily returns show little linear predictability, price volatility clusters in time—turbulent days are typically followed by turbulent days, and calm days by calm days [25, 26].
- **Asymmetric Co-movement:** Correlations between stocks tend to increase sharply during market downturns, which temporarily reduces the benefits of portfolio diversification [27, 28].

To measure co-movement without distortion from differences in individual asset volatility, each return series is normalized to zero mean and unit variance, $\tilde{r}_{i,t} = (r_{i,t} - \hat{\mu}_i)/\hat{\sigma}_i$. The resulting standardized matrix $\tilde{\mathbf{R}} \in \mathbb{R}^{T \times N}$ yields the Pearson cross-correlation matrix:

$$\mathbf{C} = \frac{1}{T-1} \tilde{\mathbf{R}}^T \tilde{\mathbf{R}}, \quad (2)$$

where each entry $C_{ij} \in [-1, 1]$ represents the pairwise linear correlation between assets i and j .

However, a fundamental practical obstacle arises when the number of assets N is comparable to the observation history T (i.e., when $Q = T/N$ is finite). Under these conditions, sample correlation matrices become heavily contaminated by random sampling fluctuations, a phenomenon known as *noise dressing* [3]. Purely accidental historical coincidences can easily masquerade as real economic relationships. Directly relying on unfiltered sample correlations can severely distort portfolio optimization, network reconstruction, and risk forecasting, creating the need for principled filtering methods.

2.2. Random Matrix Theory and Spectral Noise Filtering

Random Matrix Theory (RMT), originally introduced in mathematical physics to model the energy spectra of complex nuclei [29], provides an analytical benchmark to separate genuine economic structure from random sampling noise in empirical correlation matrices [3, 4].

Under the null hypothesis of a purely random Wishart ensemble, $\mathbf{X} \in \mathbb{R}^{T \times N}$ consists of mutually independent standard normal variables ($X_{ti} \sim \mathcal{N}(0, 1)$). In the asymptotic limit $N, T \rightarrow \infty$ with $Q = T/N \geq 1$, the eigenvalue density $\rho(\lambda) = \frac{1}{N} \sum_{k=1}^N \delta(\lambda - \lambda_k)$ of the sample covariance matrix $\mathbf{W} = \frac{1}{T} \mathbf{X}^T \mathbf{X}$ converges almost surely to the Marčenko–Pastur (MP) probability distribution [13]:

$$\rho_{\text{MP}}(\lambda) = \frac{Q}{2\pi\sigma^2\lambda} \sqrt{(\lambda_+ - \lambda)(\lambda - \lambda_-)}, \quad \text{for } \lambda \in [\lambda_-, \lambda_+], \quad (3)$$

with theoretical support bounded by $\lambda_{\pm} = \sigma^2 \left(1 \pm \sqrt{1/Q} \right)^2$, where σ^2 is the residual variance parameter.

When evaluated against empirical financial correlation matrices, the eigenspectrum $\mathbf{C} = \sum_{k=1}^N \lambda_k \mathbf{v}_k \mathbf{v}_k^T$ ($\lambda_1 \geq \lambda_2 \geq \dots \geq \lambda_N$) exhibits a characteristic tripartite decomposition [4, 5]:

1. **Market Mode** ($\lambda_1 \gg \lambda_+$): The largest eigenvalue sharply exceeds the MP upper bound λ_+ , absorbing a substantial fraction of the total matrix trace ($\text{Tr}(\mathbf{C}) = N$). Its eigenvector $\mathbf{v}_1$ has strictly positive and relatively uniform loadings ($v_{1,i} \approx 1/\sqrt{N}$), representing pervasive macroeconomic co-movement that pulls all equities simultaneously.
2. **Group/Sectoral Modes** ($\lambda_k > \lambda_+$ for $k = 2, \dots, K_{\text{signal}}$): Intermediate eigenvalues that remain above λ_+ carry genuine collective economic signal, capturing localized industry sectors, business conglomerates or shared factor exposures.
3. **Bulk Noise Component** ($\lambda_k \in [\lambda_-, \lambda_+]$): The bulk of intermediate and small eigenvalues fall within the theoretical MP band, indicating sampling noise statistically indistinguishable from a random Wishart matrix.

To isolate localized group dependencies free from market-wide drift and bulk noise, RMT enables spectral matrix reconstruction. Removing the dominant market mode ($k = 1$) while retaining genuine group modes ($k = 2, \dots, K_{\text{signal}}$) produces the group-mode filtered matrix:

$$\mathbf{C}_{\text{group}} = \sum_{k=2}^{K_{\text{signal}}} \lambda_k \mathbf{v}_k \mathbf{v}_k^T + \bar{\lambda}_{\text{rem}} \mathbf{I}, \quad (4)$$

where $\bar{\lambda}_{\text{rem}} = \frac{1}{N - K_{\text{signal}} + 1} \left(N - \lambda_1 - \sum_{k=2}^{K_{\text{signal}}} \lambda_k \right)$ preserves the trace constraint, with diagonal elements subsequently normalized such that $(\mathbf{C}_{\text{group}})_{ii} = 1$.

2.3. Topological Filtering and Planar Network Representations

The correlation matrix records the relationship between every pair of assets. This complete representation is useful for calculation, but it is difficult to interpret because all assets are connected to one another. Even after RMT filtering removes noise-compatible components, the matrix does not directly show which dependencies form the market backbone, which assets behave as hubs, or whether sectoral groups remain visible. Topological filtering addresses this problem by retaining an informative subset of the strongest relationships and representing them as a network [16].

In this network, each asset is a node and each retained dependency is an edge. Correlations are first converted into distances using the transformation proposed by Mantegna [14]: assets with stronger positive correlations are placed closer together, whereas weakly or negatively correlated assets are farther apart. This distance representation allows standard graph algorithms to select the most relevant connections. The network is therefore a simplified view of the correlation structure, not a replacement for the underlying matrix and not evidence of a causal relationship between companies.

In this study, networks are constructed from both the original correlation matrix and the RMT group-mode matrix. Comparing these representations shows which connections are mainly driven by the broad market movement and which localized structures remain after that common component is removed.

2.3.1. Minimum Spanning Tree (MST)

The Minimum Spanning Tree (MST) provides the sparsest connected representation of the market. It links every asset while preferentially retaining the shortest distances, and therefore the strongest correlations, without allowing closed loops [14, 30]. The result can be read as a market backbone: branches indicate groups of closely related assets, central nodes indicate potential hubs, and peripheral nodes represent assets that are less integrated with the rest of the network.

The MST is useful because it makes the main dependency hierarchy easy to visualize. Its simplicity is also a limitation, since alternative routes and triangular relationships are discarded. We use the MST to compare the backbone obtained from the original correlations with the backbone that remains after market-mode removal, focusing on changes in hubs and sectoral organization.

2.3.2. Planar Maximally Filtered Graph (PMFG)

The Planar Maximally Filtered Graph (PMFG) retains more of the correlation structure than the MST. Connections are considered from strongest to weakest and are added while the graph can still be represented on a plane without edge crossings [15]. Unlike the MST, the PMFG allows loops and interconnected neighborhoods. It can therefore reveal local clusters, alternative dependency paths and groups of assets that share several strong relationships.

This richer structure is particularly useful for examining whether companies from the same sector remain connected, whether a small number of firms concentrate many relevant links, and how local organization changes after RMT filtering [18, 17]. In this work, the PMFG complements the MST by supporting comparisons of clustering, sector retention, hub rankings and subsector-level dependencies. Detailed construction rules and mathematical properties are available in the original methodological references [15, 16].

2.3.3. Network Descriptors Used in This Study

Once the filtered networks are constructed, standard descriptors summarize the role of each asset and the organization of the market:

- **Degree centrality** counts how many direct connections an asset has and helps identify highly connected hubs.
- **Betweenness centrality** indicates how often an asset lies between other assets, highlighting potential bridges between groups.
- **Closeness centrality** measures how near an asset is to the rest of the network through the available paths.
- **Eigenvector centrality** gives greater importance to assets connected to other influential assets.
- **Local clustering** measures how strongly an asset's neighbors are connected to one another and helps identify cohesive local groups.

These descriptors do not represent predictive signals by themselves. They summarize structural roles in the filtered market and allow the original and group-mode networks to be compared consistently. In the retrospective volatility experiment, selected descriptors are subsequently added to conventional return and volatility variables to assess whether market position is associated with model performance. Because they are estimated from the full synchronized sample, that use remains an ex-post structural assessment rather than a real-time forecasting strategy.

2.4. Realized Volatility and Supervised Predictive Modeling

Volatility describes the intensity of price fluctuations, regardless of whether prices move upward or downward. A low-volatility period is relatively calm, whereas a high-volatility period contains larger and less stable price movements. This information is important for risk measurement, portfolio allocation and position sizing because two assets with similar average returns may expose investors to very different levels of uncertainty.

Future volatility is not known when a forecast is made. After the forecast horizon has passed, it can be approximated from the returns observed during that period. For a horizon of h trading days, the realized-volatility proxy used as a forecasting target is

$$RV_{i,t}^{(h)} = \sqrt{\sum_{\tau=1}^h r_{i,t+\tau}^2}, \quad (5)$$

where the square root returns the measure to the scale of the returns; an annualization factor can be applied when comparisons across sampling frequencies require it. This ex-post measure is called *realized volatility* [31]. In this study, the forecasting task is not to predict the direction of the next price movement. It is to estimate how much PETR4, VALE3 and BBDC4 are likely to fluctuate over the following 5 or 20 trading days.

2.4.1. How Supervised Learning Is Used

A supervised model learns from historical examples that contain two parts: information available at a given date and the volatility subsequently observed over the forecast horizon. During training, the model learns how the predictors are associated with later volatility. It then applies the learned relationship to dates that were not used for model fitting.

The conventional predictors summarize recent returns and volatility persistence. Additional feature sets introduce broad market information, RMT descriptors and the MST/PMFG network descriptors discussed in Section 2.3. The comparison includes a linear Ridge model [32], tree-based models such as Random Forest and gradient boosting [33, 34], and simple neural models [35, 36]. These approaches differ in how they combine the predictors: the linear model provides a simple baseline, tree-based models can capture nonlinear interactions, and neural models provide a limited deep-learning comparison. The purpose is to test whether adding structural information is associated with better volatility estimates, not to claim that one model represents the true dynamics of the market.

2.4.2. How Forecasts Are Evaluated

No single metric provides a complete assessment, so the models are compared using complementary measures. For M evaluation dates, the mean absolute error (MAE) and root mean squared error (RMSE) are:

$$\text{MAE} = \frac{1}{M} \sum_{t=1}^M |y_t - \hat{y}_t|, \quad \text{RMSE} = \sqrt{\frac{1}{M} \sum_{t=1}^M (y_t - \hat{y}_t)^2}. \quad (6)$$

MAE weights all absolute deviations equally, whereas RMSE gives greater weight to large errors. The out-of-sample coefficient of determination compares the model with a training-mean benchmark:

$$R_{\text{OOS}}^2 = 1 - \frac{\sum_{t=1}^M (y_t - \hat{y}_t)^2}{\sum_{t=1}^M (y_t - \bar{y}_{\text{train}})^2}. \quad (7)$$

For positive realized and predicted volatility (or variance) proxies, the Gaussian quasi-likelihood loss (QLIKE) is

$$\text{QLIKE}(y, \hat{y}) = \frac{y}{\hat{y}} - \ln \left(\frac{y}{\hat{y}} \right) - 1, \quad (8)$$

which remains informative when realized volatility is an imperfect proxy for latent market risk [37]. Lower MAE, RMSE and QLIKE values indicate better forecasts, while a higher out-of-sample R^2 indicates a stronger improvement over the benchmark.

Model configurations and feature sets are chosen with the validation period and are then evaluated on a separate chronological test period. This separation reduces the risk of selecting a model simply because it fits the evaluation sample well. The exact target construction, temporal split, model settings and statistical comparisons are reported in the methodology and results sections.

2.4.3. Retrospective Assessment vs. Real-Time Forecasting

A real-time forecast made on a particular date may use only information that was available up to that date. The return and volatility predictors satisfy this principle because they are calculated from past observations. However, the RMT and network descriptors used in this study were estimated from the complete synchronized sample. For an early date in the sample, those descriptors therefore summarize information that would not yet have been available.

Consequently, the experiment should be interpreted as a retrospective structural assessment. It tests whether the ex-post market structure is associated with volatility-model performance, but it does not establish that the same information could have improved forecasts in real time. A deployable forecasting experiment would need to reconstruct the RMT and network descriptors repeatedly, using only data available before each prediction date.

3. Related work

This section locates the manuscript at the intersection of financial dependence, RMT, filtered networks and volatility modelling. It distinguishes established tools from the contribution of the present study: an integrated B3 workflow and a clearly retrospective assessment of whether full-sample structural descriptors are associated with volatility-model performance.

3.1. Financial dependencies and econophysics

Econophysics studies financial markets using concepts and methods from statistical physics, nonlinear dynamics, complex systems and network science [38, 39]. A central observation is that financial returns exhibit stylized facts that are difficult to reconcile with independent Gaussian models. Heavy tails, volatility clustering, slow decay in absolute-return autocorrelations, intermittent bursts of activity and collective co-movement patterns are recurrent findings across asset classes, countries and sampling frequencies [1, 2]. These regularities motivate methods that focus not only on the marginal distribution of returns, but also on dependency, synchronization and network structure.

Financial dependency analysis lies at the intersection of econophysics and financial economics. In financial economics, dependence is central to portfolio risk, diversification, contagion, factor structure and systemic vulnerability. In econophysics, the same dependence is often represented through correlation matrices, eigenspectra, filtered graphs

and hierarchical structures. Network-based studies of stock markets have used this representation to connect correlation structure with market organization and temporal evolution [40, 41]. Raddant and Di Matteo [42] emphasize that correlations, causality measures, volatility spillovers and networks offer additional views of financial systems. We follow this perspective: the goal is not merely to estimate pairwise correlations, but to decompose the correlation matrix, filter its noisy components, represent the remaining structure as networks and evaluate whether the extracted network structure has forecasting relevance.

Cross-correlation dynamics during crises have received attention because market stress often increases synchronization across assets. Zheng et al. [8] showed that changes in cross-correlation structure can serve as early indicators of systemic instability. Billio et al. [6] proposed econometric measures of connectedness and systemic risk in financial sectors. Junior and Franca [7] documented that correlations among global stock indices tend to increase during crises, reducing the effective dimensionality of international markets. These findings motivate the rolling-correlation and crisis-synchronization analysis undertaken in this paper, where market-wide co-movement is treated as an object of study instead of a nuisance.

The relevance of these studies for the present manuscript is twofold. First, they show that financial dependence is dynamic and crisis-sensitive. Second, they indicate that a single unfiltered correlation matrix may conceal different layers of structure: a broad market component, group or sectoral organization and residual noise. This motivates the use of Random Matrix Theory and filtered networks as additional tools for separating and interpreting these layers.

3.2. Random Matrix Theory in financial correlations

RMT-based finance was established as an empirical programme for deciding which parts of a large correlation matrix are reproducible rather than as a new asset-pricing theory. Laloux et al. [3, 43] reported that most eigenvalues of equity correlation matrices are compatible with a random benchmark and showed that removing this component can improve portfolio construction. Plerou et al. [4] added an interpretation of the outlying eigenvectors: the leading one is a broad market mode, whereas some subleading modes are associated with sectors or other collective groups. These findings made spectral structure an empirical object in its own right, but they do not imply that every eigenvalue outside the random band has a unique economic meaning.

The methodological literature subsequently focused on how aggressively the noisy component should be treated. Bun et al. [5] review eigenvalue clipping, shrinkage and rotationally invariant estimators, and stress that the choice of cleaning rule affects both the stability of the estimated covariance matrix and its downstream use. The broader mathematical foundations are reviewed by Guhr et al. [29]. Thus, RMT studies differ in their objective: some seek a portfolio-optimal covariance estimator, while others use eigenvectors to identify collective market factors.

Our use is narrower and explicitly comparative. We use the Marčenko–Pastur edge to organize the empirical spectrum, remove the dominant market mode, retain candidate group modes and preserve the remaining trace in a residual component. The resulting matrix is not presented as an optimal covariance estimator. It is one controlled input to the paired MST/PMFG comparison, allowing us to ask whether localized topology and volatility-model performance change after the common market component is removed. Because the decomposition is estimated from the complete synchronized panel, the comparison is retrospective.

3.3. Financial networks: MST, PMFG and filtered graphs

Mantegna's introduction of the MST made it possible to study market organization without imposing a fixed correlation threshold [14]. Subsequent applications used the tree to describe sectoral and geographical taxonomies, portfolio organization and time-varying market integration. Onnela et al. [44] documented the evolution of market correlations and the changing concentration of the tree, while Bonanno et al. [30] and Song et al. [45] showed that the resulting hierarchies depend on both sector and market regime. The common result is interpretability, accompanied by a loss of cycles and alternative routes.

The PMFG line of work addresses that loss of local structure. Tumminello et al. [15, 16] used planar filtering to retain stronger links together with triangular relationships, making clustering and clique structure available for financial interpretation. Massara et al. [46] proposed the TMFG as a related scalable construction for larger systems, while temporal-network formulations retain information about the evolution of market links [47]. Other studies connected filtered topology to economic questions: Pozzi et al. [17] related peripheral positions to diversification, Aste et al. [18] examined network changes in volatile periods, and Kenett et al. [48] used partial correlations to distinguish direct from market-mediated links.

The gap relevant here is not the absence of filtering methods, but the limited integration of spectral and topological comparisons. Most network studies construct one graph from the observed correlations and interpret its hubs or communities. We instead construct matched MST and PMFG representations before and after market-mode removal, hold the asset universe and ranking procedure fixed, and compare centrality, clustering and sector retention. This design treats topology as a diagnostic of spectral filtering and as a retrospective descriptor for the forecasting experiment, rather than as evidence of causal links.

3.4. Financial networks in emerging markets

Evidence from emerging markets shows why results from large developed-market datasets cannot simply be transferred to B3. Institutional structure, sector concentration, liquidity and exposure to commodities or exchange rates can alter both the level and the persistence of cross-sectional dependence. Existing work also uses several different network objects. Tabak et al. [9] studied the Brazilian interbank network and its vulnerability to targeted failures, while Silva et al. [49] examined the position and dynamics of emerging economies in the global financial network. Wang et al. [50] compared Pearson and partial-correlation networks across world markets and found that the inferred structure depends on the dependence measure.

Studies closer to the equity-market setting remain fragmented. Pereira et al. [10] used detrended cross-correlation networks for a set of stock markets that includes Brazil; Ferreira et al. [51] compared detrended dependencies between oil and stock markets; Leonel Caetano and Yoneyama [12] analysed the Brazilian stock market with complex-network measures; and Stosić et al. [11] focused on Brazilian agrarian commodities and stocks. These contributions establish that Brazilian data contain nontrivial multiscale and network structure, but they do not jointly evaluate a broad B3 equity universe, RMT group-mode filtering, planar graphs and volatility forecasting.

Our contribution is therefore contextual rather than a claim that B3 is representative of all emerging markets. We use the original and RMT-filtered networks to separate broad market integration from more localized organization, and we report the comparison descriptively. The subsequent volatility exercise asks whether these ex-post structural differences are associated with model performance; it does not infer a universal emerging-market mechanism or a tradable signal.

3.5. Volatility forecasting and machine learning

The forecasting literature provides two relevant benchmarks. Andersen et al. [31] established the empirical role of realized volatility, while Poon and Granger [52] reviewed the comparative performance of volatility models. The HAR formulation of Corsi [53] is particularly relevant because it gives a transparent multi-horizon benchmark against which additional information can be judged. Forecast comparisons should be evaluated against strong conventional benchmarks rather than only against one another [54]. These studies motivate our use of conventional return and volatility predictors as the baseline rather than treating a flexible learner as the default.

Machine-learning work has expanded the feature space and the class of admissible nonlinearities. Bucci [35] used neural networks for realized-volatility forecasting, Christensen et al. [36] developed a broader machine-learning framework, and Zhang and Broadstock [55] connected financial-network topology with market-risk modelling. Random forests, boosting and regularized linear models remain useful comparison points because they expose different bias–variance and interpretability trade-offs [33, 34, 32].

Two methodological issues are central in this literature. First, realized volatility is a noisy proxy for latent variance, so Patton [37] recommends loss functions such as QLIKE that are robust to measurement error. Second, flexible models can produce spurious improvements when temporal ordering, feature selection or the information set is not respected. Our experiment addresses these issues with a chronological train–validation–test split, complementary losses and nested feature sets. The network variables are treated as an incremental structural augmentation: because Sets B and C use full-sample RMT and network descriptors, their results are interpreted as retrospective associations, not point-in-time forecast gains.

4. Methods: Proposed Reproducible Integrated Framework

This section specifies the proposed reproducible integrated framework. Its contribution is methodological integration rather than a new RMT, MST or PMFG algorithm: four established layers that are often studied separately are connected in one auditable sequence—(i) B3 data and asset-universe construction, (ii) spectral separation of market, group and noise components, (iii) matched topological representations of original and filtered dependencies, and (iv)

clearly labelled retrospective model assessments. The same sequence can be reapplied to another market or sample by replacing the data and asset labels while preserving the decision rules. Figure 1 summarizes the framework.




Figure 1: Overview of the workflow. The workflow proceeds through stylized-fact checks, correlation estimation, Random Matrix Theory filtering, hierarchical clustering, MST/PMFG network construction, subsector aggregation and volatility forecasting with machine-learning, deep-learning, RMT and network-based features.

4.1. Data and asset universe

This first block defines the empirical universe before any dependence is estimated. It separates raw historical coverage, the complete panel required by spectral methods and the smaller set of representative assets used for the volatility experiment.

4.1.1. BovDB/B3 data source and adjusted closing prices

We use daily Brazilian equity prices from BovDB and BovDBv2 [56, 57]. BovDB provides an open historical B3 price database, while the public BovDBv2 release extends the coverage and also contains intraday records. We retain daily adjusted closing prices because the objective is long-horizon cross-sectional analysis; daily sampling provides a common frequency across assets and limits the microstructure effects that are more prominent in intraday data. Adjustments for splits, dividends, interest on equity, bonuses and related corporate actions are essential because unadjusted jumps would contaminate returns, correlations and graph weights. The local daily-candle extraction used for the descriptive 1998–2025 coverage is distinguished from the public 2024 BovDBv2 release and is not redistributed with the software repository.

4.1.2. Asset filtering and universe definition

We apply deterministic filters to retain direct listed equities with valid, strictly positive adjusted prices. BDRs, ETFs, investment funds, indices and other non-equity vehicles are excluded. We inspect extreme adjusted-return observations as a quality-control step and remove an asset only when the evidence indicates an implausible corporate-action or ticker-history inconsistency. A minimum historical-coverage rule is then applied, one share class per issuer is retained where necessary, and the remaining assets are synchronized on common trading dates.

The resulting core historical universe contains 58 equities. Its raw unbalanced file spans 17 March 1998 to 30 December 2025, whereas the complete synchronized panel used for RMT, clustering and networks contains $N = 58$ assets and $T = 1527$ common observations from 23 August 2006 to 15 August 2024. The distinction prevents pairwise descriptive samples from being confused with the common panel required by the spectral analysis.

Core historical universe. The universe is historical rather than restricted to currently listed firms. This preserves a long B3 cross-section while keeping the number of assets compatible with interpretable spectral and network visualizations. Table 1 distinguishes the broader representative-asset coverage from the complete synchronized panel used for RMT, clustering and network construction.

Table 1: Summary of the samples used in this study.

Sample	Assets	Period	Main use
Representative assets	3	2006–2025	Stylized facts and initial forecasting experiment
Core historical universe	58	2006-08-23–2024-08-15	RMT, clustering and networks (complete panel)

Display convention: The global numerical-precision convention in Section 4.4 applies.

The complete-case restriction is important because the empirical correlation matrix must be estimated from a common set of trading dates across all assets. Static pairwise and rolling correlations may use the observations available for each pair or window; they therefore retain broader unbalanced coverage and should not be interpreted as estimates from the complete RMT panel.

Sector and subsector classification. Each asset in the core universe is assigned to a macro-sector, subsector and segment classification. The macro sectors include Financials, Oil Gas and Biofuels, Basic Materials, Utilities, Industrials, Consumer Cyclical, Consumer Non-Cyclical, Real Estate and Telecommunications. These labels are used for sectoral correlation comparisons, heatmap interpretation, sector-retention checks and subsector-level aggregation. Utilities and Real Estate are kept separate as a descriptive classification, not as a claim that all firms in either group share the same risk profile. Within Basic Materials, the subsector map distinguishes Mining, Steel and Metallurgy, Chemicals, and Paper and Cellulose; detailed asset assignments are provided in the supplementary material.

Data quality notes and limitations. The main quality concern is the distinction between genuine extreme returns and artificial jumps caused by corporate-action adjustment problems. Adjusted prices reduce the impact of splits, dividends and related events, but do not guarantee that every historical corporate action, ticker transition, share-class change or reorganization is perfectly corrected. Extreme-return validation is therefore treated as a separate quality-control step, and assets with implausible corporate-action-related jumps are excluded from the main analyses. The complete synchronized panel provides a consistent input for RMT and networks but shortens the available historical window; results may differ under intraday sampling, alternative liquidity filters, rolling universes or broader forecasting targets.

4.1.3. Log returns and synchronized panel

Daily log returns follow the definition in Section 2.1, Eq. (1), and are applied to the adjusted closing-price series retained after filtering. The resulting synchronized return panel contains the same $N = 58$ assets on the $T = 1527$ common dates. Complete-case synchronization is performed before any dependence estimation so that eigenvalues, eigenvectors, distances and network edges refer to one fixed date–asset panel.

4.1.4. Stylized facts and correlation diagnostics

The stylized-fact interpretation and standardization conventions are those defined in Section 2.1. PETR4, VALE3 and BBDC4 are used as liquid, economically distinct demonstration assets for return diagnostics and as initial forecasting targets. We report return distributions, tail summaries, moments, empirical Value at Risk and autocorrelations of absolute returns. The Pearson correlation matrix is constructed according to Eq. (2); this subsection specifies only its

empirical use. Within- versus between-sector correlations are compared using the B3 classifications and an asset-level sector-label permutation check. Rolling average pairwise correlation is used as a descriptive synchronization measure. These diagnostics establish the empirical object passed to the spectral and network stages; they are not treated as trading signals or causal tests.

4.2. RMT spectral decomposition and mode separation

The spectral step follows the RMT formulation in Section 2.2, including the Marčenko–Pastur density and support in Eq. (3) and the trace-preserving group-mode reconstruction in Eq. (4). Operationally, we apply that rule to the complete synchronized correlation matrix, order the eigenvalues from largest to smallest, remove the leading market mode, retain subsequent eigenmodes whose eigenvalues exceed the upper MP edge as candidate group components, and represent the remaining components by the residual term. The reconstructed matrix is normalized as a correlation matrix before network construction. This is a filtering convention rather than a claim that the MP null fully describes the Brazilian market; the “group” label denotes spectral separation, not a unique causal factor. Issuer-collapsed and circular-shift analyses are used as robustness checks for this descriptive interpretation.

4.3. Filtered networks and hierarchical representations

The same distance transformation and edge-ranking procedure is applied to the original correlation matrix and to the RMT group-mode matrix. This matched construction is central to the proposed methodology: differences in topology can then be attributed to the spectral input rather than to a change in the asset universe or graph algorithm.

4.3.1. MST, PMFG and hierarchical clustering

We use the Mantegna distance $d_{ij} = \sqrt{2(1 - C_{ij})}$, average-linkage clustering and ordered heatmaps to inspect hierarchical organization. The MST provides the sparse market backbone, while the PMFG retains additional planar links and local clustering. For each representation we compare degree, betweenness, closeness, eigenvector centrality and local clustering, together with sector retention and hub reorganization. The mathematical definitions of these standard descriptors are given in Section 2.3; here they are operationalized as outputs of the paired original/group-mode construction. Cophenetic agreement is reported using the standard dendrogram-comparison perspective of Sokal and Rohlf [58].

4.3.2. Subsector dependency aggregation

Following the parallel MST/PMFG and hierarchical-clustering views, ticker-level dependency information is summarized as a mesoscopic economic representation. Retained ticker-level edges are assigned to pairs of B3 subsectors. For subsectors a and b , the displayed weight is the mean retained dependency,

$$W_{ab} = \frac{1}{|\mathcal{E}_{ab}|} \sum_{(i,j) \in \mathcal{E}_{ab}} w_{ij}, \quad \mathcal{E}_{ab} = \{(i,j) \in E : g(i) = a, g(j) = b\}. \quad (9)$$

This operation compresses many asset-level links into an interpretable subsector-to-subsector map. It is used to identify whether filtered dependencies concentrate within or between subsectors and to highlight bridge-like groups. The result is a descriptive aggregation of the filtered graph, not a second correlation estimate or a causal channel model.

4.4. Retrospective model assessments

This final block asks whether the structural descriptors produced above are associated with model performance. For $i \in \{\text{PETR4, VALE3, BBDC4}\}$ and $h \in \{5, 20\}$, the target is the forward realized-volatility measure defined in Eq. (5). Training covers 2006–2017, validation 2018–2020 and the held-out test period 2021–2025; this chronological ordering is used for model fitting, feature-set selection and final evaluation. Only variables available through date t can support a genuine real-time forecast: Set A satisfies this requirement, whereas Sets B and C use full-sample structural descriptors and are therefore interpreted as retrospective structural associations. A deployable forecasting claim would require reconstructing RMT and network features recursively at every forecast origin.

The nested feature sets provide the operational comparison. Set A contains lagged return and volatility variables, Set B adds rolling market variables and full-sample RMT descriptors, and Set C adds full-sample MST/PMFG descriptors. Ridge, Random Forest and histogram gradient boosting are evaluated alongside compact MLP and CNN-1D models. The fixed feature sets and model menu test incremental association under a common design; they are not intended as a causal feature test or an exhaustive architecture search.

Forecast quality is reported using the four metrics defined in Section 2.4: MAE, RMSE, QLIKE and R^2_{OOS} . Lower MAE, RMSE and QLIKE indicate better forecasts, while higher R^2_{OOS} indicates improvement over the benchmark. Fixed-seed circular moving-block bootstrap intervals summarize test uncertainty, and a one-sided HAC comparison provides a focused retrospective contrast between the validation-selected structural and Set A specifications. The global numerical-display convention is applied to all reported counts, estimates, metrics, intervals, eigenvalues, graph quantities and test statistics.

5. Stylized Facts and Correlation Structure

This section establishes the empirical baseline for the subsequent spectral and network analyses. We first describe return regularities in the three demonstration assets, then summarize the cross-sectional correlation distribution, compare sectoral groups and track broad synchronization over time. The results are descriptive: they show the dependency patterns that later RMT and graph representations are intended to organize.

5.1. Stylized facts of selected B3 assets

Figure 2 summarises the return properties of PETR4, VALE3 and BBDC4. We use these assets as representative examples and as the initial forecasting targets because they are liquid, economically important and sectorally distinct; they are not intended to be a proxy for the entire Brazilian equity market. The figure combines four checks that are standard in the econophysics literature [2]: normalized prices, daily returns, the complementary cumulative distribution of absolute returns and the autocorrelation of absolute returns.




Figure 2: Stylized facts for PETR4, VALE3 and BBDC4 over 2006–2025. The panels show normalized prices, daily log returns, the complementary cumulative distribution of absolute returns and autocorrelations of absolute returns. The return panel uses three vertically aligned subpanels with a common scale and a zero reference line for each asset.

The normalized-price panel illustrates heterogeneous long-run trajectories across the three assets. This heterogeneity is expected because the selected firms belong to distinct economic segments and are exposed to different sectoral and macro-financial drivers. The return panel displays periods in which large movements cluster over time, a pattern

indicating volatility clustering. The CCDF panel suggests that large absolute returns occur more frequently than would be expected under a Gaussian benchmark, pointing to heavy-tailed returns. The autocorrelation panel shows positive dependence in $|r_t|$ over several lags, indicating persistent volatility [2].

5.2. Descriptive statistics

Table 2 quantifies the same diagnostics. All three series have means close to zero, substantial kurtosis and non-negligible counts of large absolute returns. PETR4 has the largest negative daily observation in the 2006–2025 window, while VALE3 and BBDC4 also show stress-period extremes. The autocorrelations of absolute returns remain positive across several lags, especially for VALE3.

The descriptive statistics match the visual reading of Figure 2. The return distributions display substantial departures from a Gaussian benchmark, including excess kurtosis and persistent absolute-return autocorrelations. These properties match the well-documented stylized facts of financial returns [2] and support treating the Brazilian market as a complex system in which dependency structure, not marginal normality, is the relevant modelling object.

Table 2: Descriptive statistics for the three demonstration assets, 2006–2025.


Symbol	n	Mean	Std.	Skew.	Ex. kurt.	Min.	Max.	VaR _{1%}	VaR _{5%}	$ r > 10\%$	$ r > 20\%$	ACF ₁	ACF ₅	ACF ₂₀	ACF ₆₀
BBDC4	4953	0.0003	0.0218	0.02	7.96	-0.1910	0.1999	-0.0559	-0.0312	16	0	0.207	0.174	0.122	0.044
PETR4	4953	0.0004	0.0276	-0.76	11.55	-0.3524	0.2007	-0.0767	-0.0408	41	4	0.254	0.209	0.140	0.062
VALE3	4953	0.0005	0.0258	-0.22	7.51	-0.2814	0.1936	-0.0676	-0.0381	26	2	0.231	0.207	0.169	0.119

Notes: Mean, Std., extrema and VaR are reported to four decimal places; skewness and excess kurtosis to two; autocorrelations to three; and event counts as integers. Mean and Std. are computed from daily log returns. VaR_{1%} and VaR_{5%} are empirical lower-tail return quantiles, and ACF _{ℓ} is the autocorrelation of absolute returns at lag ℓ trading days.

Display convention: The global numerical-precision convention in Section 4.4 applies.

5.3. Static correlation distribution

We now turn from individual return behaviour to cross-sectional dependence. The core historical universe contains 58 assets, producing $N(N - 1)/2 = 1653$ unique asset pairs. The static pairwise summaries below use the available overlap for each pair and therefore describe the broader unbalanced historical coverage; Figure 3 summarizes their distribution and then splits pairs into within-sector and between-sector groups. The subsequent RMT decomposition instead uses the single complete 58-asset panel from 23 August 2006 to 15 August 2024.



(a) Full pairwise distribution.

(b) Within-sector and between-sector pairs.

Figure 3: Correlation structure in the core historical universe. Panel (a) shows the full Pearson correlation distribution. Panel (b) compares within-sector and between-sector correlations using the initial macro-sector classification.

The full distribution is centred in a moderate positive range, with mean correlation close to 0.24 and median correlation close to 0.23. This level of average dependence suggests the presence of a common component, but the distribution remains dispersed. The dispersion is important because it leaves room for sectoral and subsectoral organization instead of a single homogeneous market block.

5.4. Sectoral correlation comparison

The sector split gives a descriptive check of this interpretation. Within-sector correlations are visibly shifted to the right relative to between-sector correlations. Table 3 reports the corresponding pair counts and summary statistics. These results indicate that the initial macro-sector map is reflected in the dependency matrix.

Table 3: Within-sector and between-sector correlation summary.

Group	Pairs	Mean	Median	Std. dev.	Min	Max
Within-sector	275	0.3433	0.3082	0.1819	0.0534	0.9402
Between-sector	1378	0.2249	0.2076	0.1166	-0.0236	0.5552

Display convention: The global numerical-precision convention in Section 4.4 applies.

The difference is economically meaningful. Within-sector pairs have mean correlation 0.3433, while between-sector pairs have mean correlation 0.2249, giving $\Delta_{\text{sector}} = 0.1184$. An asset-level label-permutation test with 10,000 sector-label reallocations preserves both the observed dependency matrix and the sector-size vector, without treating overlapping pairs as independent observations. None of the permuted statistics reaches the observed difference (null 95th percentile = 0.0245; one-sided Monte Carlo $p < 0.0001$). Conditional on exchangeability of sector labels under the null and on the initial macro-sector map, this provides evidence that the sectoral grouping is reflected in the dependency matrix.

5.5. Dynamic market synchronization

Full-sample correlations are limited because market dependence changes over time. Figure 4 shows the rolling average market correlation using 252- and 504-trading-day windows. This statistic compresses the correlation matrix into a time series that measures how synchronized the market is at each point in time. The historical panel has changing coverage in its earliest years: each window uses the assets with data available in that window, ranging from 40 assets initially to the full 58-asset universe in later periods. The series should therefore be interpreted as a descriptive measure of the average available pairwise dependence.




Figure 4: Rolling average market correlation for the core historical universe, estimated with 252- and 504-trading-day windows. Each window uses the assets available during that period.

The rolling average correlation tends to rise around major stress periods, including the 2008 global financial crisis, the Brazilian recession period around 2015–2016 and the 2020 COVID shock. These episodes indicate periods

of stronger market-wide synchronization, which helps explain the dominant leading eigenvalue in the full-sample correlation matrix.

Figure 5 adds a more granular view by showing selected pairwise dynamic correlations under different estimators. The pairs were selected for their economic interpretability: PETR4–VALE3 compares two large commodity-exposed firms in distinct subsectors; BBDC4–BBAS3 represents financials; GGBR4–USIM5 represents steel and metallurgy; and ELET3–CMIG4 represents utilities. The figure is illustrative rather than a representative sample of all possible pairs.




Figure 5: Dynamic pairwise correlations for selected economically interpretable asset pairs. Pearson estimates use 20- and 60-trading-day windows; EWMA uses a 20-day half-life. The estimates show how individual dependencies can vary across time scales.

The analytical role of these estimates is to show that a single full-sample correlation can conceal episodes in which a dependency strengthens, weakens or temporarily breaks down. The 20-day Pearson estimate reacts quickly to

local changes but is noisy; the 60-day estimate is smoother but slower; and EWMA with a 20-day half-life provides an intermediate adaptive estimate. Together, the three views distinguish persistent co-movement from short-lived correlation bursts. This makes the figure a useful sensitivity check for the static correlation matrix used in the subsequent RMT and network analyses: it does not replace the full matrix, but demonstrates why its resulting structures should be interpreted as long-run summaries of a time-varying dependency system.

6. Random Matrix Filtering

This section asks which parts of the empirical correlation matrix are compatible with the RMT sampling-noise benchmark and which are candidates for collective structure. We report the eigenspectrum, reconstruct interpretable components and use robustness checks to delimit the economic reading of the modes. The results describe full-sample structure; they do not estimate time-varying factors.

6.1. Empirical eigenspectrum

Random Matrix Theory supplies the spectral benchmark for distinguishing structured collective modes from sampling noise in the empirical correlation matrix [3, 4]. The core historical matrix uses $N = 58$ assets and $T = 1527$ observations, giving $Q = T/N \approx 26.33$. The Marčenko–Pastur bounds are $\lambda_- = 0.6482$ and $\lambda_+ = 1.4278$.

Table 4: RMT summary for the core historical universe.

N	T	Q	λ_-	λ_+	λ_1	λ_2	λ_3	λ_4	λ_5	Above λ_+
58	1527	26.33	0.6482	1.4278	21.6505	3.0957	1.7839	1.6543	1.4835	5

Display convention: The global numerical-precision convention in Section 4.4 applies.

Table 4 reports the numerical inputs for RMT filtering. With $N = 58$ assets and $T = 1527$ complete observations, the aspect ratio is $Q = 26.33$, which implies a relatively narrow Marčenko–Pastur noise band from $\lambda_- = 0.6482$ to $\lambda_+ = 1.4278$. The first eigenvalue, $\lambda_1 = 21.6505$, is therefore not a marginal deviation from the random-matrix benchmark but a dominant collective component. The next four eigenvalues, $\lambda_2 = 3.0957$, $\lambda_3 = 1.7839$, $\lambda_4 = 1.6543$ and $\lambda_5 = 1.4835$, also exceed the upper random bound. There are 5 sample eigenvalues outside the band: these are the candidate collective modes, indicating one broad market mode and four additional modes that may be associated with group, sectoral or localized dependency structure.

Figure 6 shows the spectrum against the random benchmark. The largest eigenvalue lies far above the Marčenko–Pastur upper bound. Four additional eigenvalues also exceed the upper edge, indicating a dominant market mode and a small number of group or sector modes.




Figure 6: Empirical eigenvalue spectrum and Marčenko–Pastur bounds. The largest eigenvalue is far outside the random band and captures the market-mode component.

6.2. Market Mode and significant eigenvalues

The largest eigenvalue is 21.6505, corresponding to approximately 37.3% of the total trace of the correlation matrix. Its eigenvector has broadly positive loadings, confirming its interpretation as a market mode. It captures broad co-movement and explains why the original correlation matrix has a strong positive common component. The next four eigenvalues are 3.0957, 1.7839, 1.6543 and 1.4835. These values are above λ_+ and act as candidates for group or sector structure.

The market mode is meaningful but visually dominant. If a network is built directly from the original matrix, many edges can reflect market-wide co-movement instead of local structure. We construct the group-mode matrix to ask what remains after the broad common component is removed.

6.3. Matrix reconstruction

The reconstruction step checks that the spectral decomposition and the filtered matrices are being computed accurately. If all eigenvalue–eigenvector modes are recombined, the original correlation matrix is recovered up to floating-point precision. Table 5 reports this numerical check: the Frobenius reconstruction error is approximately 2.74×10^{-14} , and the maximum absolute reconstruction error is approximately 3.99×10^{-15} . The reported unit diagonal refers to the filtered matrix after the partial reconstruction has been rescaled for downstream correlation-based distances. This check reduces the possibility that subsequent network differences are driven by numerical reconstruction error rather than by the intended selection of market, group and noise-compatible modes.

Table 5: RMT reconstruction diagnostics.

Diagnostic	Value
Frobenius reconstruction error	2.74×10^{-14}
Maximum absolute reconstruction error	3.99×10^{-15}
C_{filtered} diagonal min	1.0000
C_{filtered} diagonal max	1.0000
Eigenvalues above λ_+	5
Market-mode trace share	0.3733

Display convention: The global numerical-precision convention in Section 4.4 applies.

Figure 7 contrasts the original correlation matrix with the candidate group/sector component. This focused comparison makes the effect of market-mode removal readable in the main text. The market, noise-compatible and retained components are reported in the supplementary full decomposition.

Empirically, the eigenvectors beyond the leading component show loading patterns suggesting localized or sector-related dependencies. This residual organization motivates the clustering and network sections: if localized dependency patterns remain after market-mode removal, MST and PMFG structures should differ when constructed from the group-mode matrix.




Figure 7: Original and candidate group/sector RMT components at readable scale. The original matrix is dominated by broad positive co-movement, whereas the group/sector component highlights weaker localized contrasts after the leading market mode is removed.

6.4. Top eigenvector loadings

Eigenvalues identify the strength of collective modes, but eigenvectors help interpret their composition. Figure 8 reports the largest loadings for the five modes above the Marčenko–Pastur upper bound. For each eigenvector, assets are ranked by the magnitude of their loadings. The signs of eigenvector loadings are arbitrary up to multiplication by -1 ; the relevant information is therefore the relative contrast between groups of assets, together with the magnitude and concentration of their loadings.

The first eigenvector is a broad market mode. All 58 assets have positive loadings, ranging from 0.037 to 0.174, and the leading mode accounts for 37.3% of the trace of the correlation matrix. The largest loadings include steel, mining and financial assets, but the positive contribution is distributed across all sectors. This pattern supports its interpretation as market-wide co-movement rather than a sector-specific factor.

The second eigenvector ($\lambda_2 = 3.10$) is best read as a relative contrast between mining and steel firms—including VALE3, BRAP3/4, GGBR3/4, GOAU3/4, CSNA3 and USIM3/5—and a group led by utilities and selected consumer-cyclical assets. It therefore contains a commodity-related dimension, but it is not a pure Basic Materials factor. The third eigenvector ($\lambda_3 = 1.78$) similarly contrasts the financial pole formed by BBDC3/4, ITSA4 and BBAS3 with a more heterogeneous industrial, chemical and consumer-cyclical pole led by ROMI3, SHUL4, BRKM3, ALPA4 and GUAR3.

The fourth eigenvector ($\lambda_4 = 1.65$) is highly concentrated in TELB3 and TELB4, the two listed TELEBRAS share classes. We interpret this result as an issuer-specific share-class component, rather than as evidence of a broad telecommunications factor. The fifth eigenvector ($\lambda_5 = 1.48$) separates a coherent utilities block—including CEMIG, COPEL, Eletrobras, CPFL Energia and ISA CTEEP—from financial and selected consumer-cyclical assets. This is a plausible utilities-related pattern, but λ_5 lies only about 3.7% above the Marčenko–Pastur upper bound. Its economic interpretation should therefore remain cautious and be treated as a candidate localized mode.

Overall, Figure 8 shows that the modes beyond the market component are structured contrasts, rather than one-sector classifications. They can mix sectors when firms share macroeconomic, liquidity or balance-sheet exposures. This interpretation motivates the subsequent clustering and network analyses, which test whether these spectral contrasts reappear in other representations of the dependency matrix.

6.5. Issuer-level share-class robustness

The 58-asset panel contains 12 companies represented by two listed share classes. This can amplify issuer-specific co-movement and, in the case of TELEBRAS, produces a component that is not an economically broad sector factor. We therefore repeat the RMT calculation after retaining one deterministic representative per company, reducing the panel to 46 issuers. The retained series is the class with the most valid daily returns in the raw historical file, with alphabetical tie-breaking; this rule is stated in Section 4.

Table 6: Issuer-level RMT robustness check. The collapsed panel retains one share class per company.

Panel	N	T	Q	λ_+	Modes above λ_+	λ_1/N
Full share-class panel	58	1527	26.33	1.4278	5	0.3733
Issuer-collapsed panel	46	2207	47.98	1.3096	3	0.3363

Display convention: The global numerical-precision convention in Section 4.4 applies.

The dominant market component is robust: its trace share remains 33.6% after issuer collapse, compared with 37.3% in the full panel. Three eigenvalues remain above the issuer-panel bound, so the market mode and two additional collective components are not artifacts of duplicated share classes. In contrast, the TELEBRAS share-class component necessarily disappears, and the two weaker full-panel modes do not remain above the issuer-panel comparison. The principal market-structure result is therefore robust, whereas interpretations of marginal localized modes should account for issuer representation.

RMT-filtered networks: retrospective volatility assessment




Figure 8: Top eigenvector loadings for the five RMT modes above the Marčenko–Pastur upper bound. Eigenvector 1 is a broad market mode. Subsequent eigenvectors show relative contrasts between groups of assets; their overall signs are arbitrary.

6.6. Circular-shift null robustness

The Marčenko–Pastur benchmark assumes an idealized independent-noise setting. As a complementary check, we compare the five leading eigenvalues with 1,000 independent circular-shift surrogates. Each surrogate preserves the temporal ordering and marginal distribution of every individual return series while breaking their same-date alignment. Table 7 reports rank-specific 95th-percentile thresholds and Monte Carlo exceedance probabilities.

Table 7: Circular-shift null check for the five leading eigenvalues. The null uses 1,000 independent circular shifts, one per asset and replicate.

Rank	Empirical λ	Null 95th percentile	Monte Carlo p
1	21.6505	1.6882	< 0.001
2	3.0957	1.4932	< 0.001
3	1.7839	1.3889	< 0.001
4	1.6543	1.3533	< 0.001
5	1.4835	1.3295	< 0.001

Display convention: The global numerical-precision convention in Section 4.4 applies.

All five empirical eigenvalues exceed the rank-matched 95th-percentile circular-shift threshold, including the fifth mode. Thus, its departure from the noise bulk is not explained solely by the temporal dependence or non-Gaussian marginal behavior retained by this surrogate null. This result should be read jointly with the issuer-collapsed check: the fifth mode is statistically distinguishable from this synchronization-free null, but its detailed economic interpretation remains sensitive to how multiple share classes are represented.

7. Hierarchical Clustering and Ordered Heatmaps

This section examines whether the correlation patterns identified by RMT remain visible when assets are ordered by their distance geometry. Ordered heatmaps reveal broad and localized blocks, while the dendrogram provides one readable summary of group-mode proximity. These are visualization and organization tools, not a rule that assigns every asset to a fixed economic cluster.

7.1. Ordered heatmaps

RMT decomposes the correlation matrix into spectral components, while ordered heatmaps offer a visual check of whether these components display block-like organization. We order the assets through hierarchical clustering in Figure 9, which displays original, filtered and group-mode structures. The heatmaps serve as diagnostic tools, assessing whether spectral filtering is associated with patterns that can be read in relation to sector and subsector classifications.

The original matrix displays a broad positive background, reflecting the large leading eigenvalue. The filtered matrix preserves part of this organized dependence while reducing the contribution of eigenvalues inside the random-matrix band. The group-mode matrix has lower average correlation by construction, but its ordered representation displays localized block-like patterns that are less dominated by the market-wide component.




Figure 9: Ordered heatmaps for original and RMT-filtered correlation matrices. Panel titles are enlarged and the common color scale makes the group-mode blocks directly comparable with the original and retained structures.

7.2. Dendrogram comparison

Dendrograms give another view of the same distance geometry. We show the group-mode tree in Figure 10, because it is the representation that most directly addresses the structural question after market-mode removal. The colors identify the initial macro-sector classification, making the figure a visual diagnostic rather than a sector-assignment rule. Table 8 retains the original, filtered and group-mode fit summaries; the full three-tree comparison is supplementary material.




Figure 10: Group-mode correlation-matrix dendrogram using Mantegna distances and average linkage. It summarizes the localized geometry remaining after the broad market component is removed; colored labels indicate the initial macro-sector classification.

Table 8: Cophenetic correlations for hierarchical clustering.

Matrix	Assets	Cophenetic correlation
Original	58	0.950
Filtered	58	0.934
Group Mode	58	0.871

Display convention: The global numerical-precision convention in Section 4.4 applies.

The cophenetic correlation is highest for the original matrix, 0.950, and remains high for the filtered matrix, 0.934. The group-mode value is lower, 0.871. Removing the leading market component reduces the smooth common structure

that is easiest for a tree to summarize. The lower value therefore indicates a less tree-like group-mode geometry, not an absence of localized organization.

The clustering results connect naturally to network filtering. Ordered heatmaps and dendrograms suggest that localized structure remains after RMT filtering. We next examine how this structure appears when only the strongest distance-ranked edges are retained under MST and PMFG constraints.

8. Financial Network Topology

Clustering summarises hierarchical organisation, but it does not identify graph-theoretic hubs, bridges or local cycles. Financial networks give these additional views by retaining selected high-correlation, short-distance edges under structural constraints. We move from dendrograms to filtered financial networks to examine how dependency structure changes after RMT filtering.

8.1. MST backbone

The MST is the sparsest connected representation of the dependency geometry. With $N = 58$, it contains 57 edges. Figure 11 compares the MST built from the original matrix with the MST built from the group-mode matrix. This comparison illustrates how the MST backbone changes after the leading market component is removed.




Figure 11: Refined MST comparison for original and group-mode correlation matrices. The original MST emphasizes strong market-mode dependencies, while the group-mode MST highlights weaker localized dependency channels.

The original MST has mean edge correlation 0.5796, while the group-mode MST has mean edge correlation 0.1388. This decline is expected because the group-mode network is constructed after excluding the dominant common component. The same-sector edge ratio also changes, from 0.7193 in the original MST to 0.6316 in the group-mode MST. The group-mode MST is not simply a weaker version of the original tree. It represents a different dependency backbone constructed after removing the leading common component.

8.2. PMFG topology

The PMFG is less sparse than the MST because it preserves planarity while retaining $3N - 6 = 168$ edges. It can retain cycles, triangles and 4-cliques, which are absent from the MST. Figure 12 compares original and group-mode PMFGs. This is the main network visualization in the paper because the PMFG retains enough structure to study local connectivity patterns, clustering and hub reorganization.

The original PMFG has mean edge correlation 0.5177. The group-mode PMFG has mean edge correlation 0.1109. Despite this lower average edge strength, the group-mode PMFG has higher average clustering, 0.6653, than the original PMFG, 0.5273. Removing the market mode lowers the absolute level of edge correlations, but the resulting PMFG is associated with stronger local clustering.




Figure 12: Refined PMFG comparison for original and group-mode correlation matrices. The group-mode PMFG reduces market-wide dominance and highlights more localized dependency channels.

8.3. Original versus group/sector-mode networks

Table 9 summarizes the main MST and PMFG statistics. Across both graph families, original networks display stronger average edge correlations, while group-mode networks display weaker edge correlations but retain nontrivial sectoral and graph structure. Original networks have stronger average edge correlations, as expected, because they are constructed from matrices that include the leading market component. Group-mode networks have weaker edges but retain patterns aligned with sectoral and subsectoral groupings.

Table 9: Selected MST and PMFG topology statistics.

Network	Matrix	Edges	Mean edge corr.	Same-sector ratio	Clustering	Top degree	Top betweenness	4-cliques
MST	Original	57	0.5796	0.7193	0.0000	RAPT4	ITSA4	0
MST	Group mode	57	0.1388	0.6316	0.0000	ALPA4	ALPA4	0
PMFG	Original	168	0.5177	0.6012	0.5273	BBDC4	GGBR4	55
PMFG	Group mode	168	0.1109	0.5417	0.6653	GUAR3	GUAR3	55

Display convention: The global numerical-precision convention in Section 4.4 applies.

8.4. Temporal stability of network features

The static graphs summarize the full synchronized sample, but their components need not be stable over time. We therefore re-estimate each network in three consecutive 509-trading-day windows: 23 August 2006–13 February 2012, 28 February 2012–2 June 2021, and 4 June 2021–15 August 2024. Table 10 reports the mean pairwise overlap across the three window pairs. Edge persistence is modest in the original networks (mean Jaccard 0.206 for MST and 0.212 for PMFG) and lower in the group-mode networks (0.047 and 0.140, respectively). Top-five betweenness-hub overlap is only 0.20–0.27 on average, and community ARI values are similarly low.

No top-five betweenness hub appears in all three windows; the most recurrent hubs occur in two. Thus the full-sample graph identifies long-run average dependency channels, while its individual edges, central nodes and modular partitions are sensitive to the observation window. The lower persistence of group-mode edges is consistent with a weaker component reconstructed after removal of the dominant market mode. This result strengthens the descriptive interpretation of Figures 11–14: they reveal economically interpretable configurations in the full sample, not invariant network structure.

Table 10: Temporal stability across three consecutive 509-trading-day windows. Values are means over the three window pairs.

Network	Matrix	Edge Jaccard	Top-5 hub overlap	Community ARI
MST	Original	0.206	0.267	0.138
MST	Group mode	0.047	0.200	0.182
PMFG	Original	0.212	0.200	0.214
PMFG	Group mode	0.140	0.200	0.174

Display convention: The global numerical-precision convention in Section 4.4 applies.

8.5. Network topology metrics

Figure 13 expands the comparison beyond average edge correlation. It summarizes mean edge correlation, same-sector ratio, same-subsector ratio and average clustering across network and matrix types. This figure illustrates that network structure cannot be reduced to a single edge-weight statistic.

The network comparison supports a subsector-level interpretation. Group-mode networks have lower mean edge correlations, while the PMFG clustering coefficient increases. Same-sector and same-subsector ratios remain substantial after market-mode removal. Table 11 confirms that all PMFG variants satisfy the expected planar clique counts. This pattern suggests localized dependency structure, although a formal comparison with graph null models would be required to quantify how unlikely the structure is under randomized alternatives. MST clustering is zero by construction because trees do not contain triangles; the PMFG clustering coefficient is the relevant statistic for local triangular organization.

Table 11: PMFG clique and planarity diagnostics.

Matrix	Nodes	Edges	Triangles	4-cliques	Planar
Original	58	168	166	55	Yes
Filtered	58	168	166	55	Yes
Group Mode	58	168	166	55	Yes

Display convention: The global numerical-precision convention in Section 4.4 applies.

The clique counts are construction checks. For a PMFG with 58 nodes, 168 edges are expected. The presence of triangles and 4-cliques shows that the PMFG retains local graph structures that are absent from the MST. These structures can be transformed into network predictors for the forecasting experiment, although their economic interpretation depends on the underlying edge weights and sector composition.

8.6. Hub reorganization after RMT filtering

Hub rankings summarize which assets occupy central positions in the constructed networks. Figure 14 compares hub rankings before and after filtering. In the original PMFG, GGBR4 is the top betweenness node and BBDC4 is the top degree node. In the group-mode PMFG, GUAR3 becomes the leading node by both degree and betweenness.

This change in hub ranking is meaningful for interpretation. Original hubs can reflect market-wide correlation strength and broad exposure to the common component. Group-mode hubs act as localized bridges in the network obtained after removing the leading common component. The appearance of assets such as GUAR3, BBDC3, VALE3, CCRO3, SAPR4 and CMIG3 in group-mode rankings suggests that centrality depends on the matrix used to construct the network. Because hub identities can be sensitive to small changes in edge ranking, we report these rankings as descriptive summaries rather than stable structural labels.




Figure 13: Network structure comparison across original, filtered and group-mode MST/PMFG representations.


Hub Shift: Impact of RMT Filtering on Network Centrality

Figure 14: Hub-rank comparison across original and group-mode networks. RMT filtering changes which assets occupy central positions in the dependency network.

9. Aggregated Subsector Dependency Network

This final structural summary translates retained ticker-level edges into a subsector-level map: nodes are subsectors and an edge is the average retained dependency between their constituent tickers. It is a descriptive aid for identifying concentrated or bridge-like connections, not a causal map or a new aggregation estimator. Figure 15 and Table 12 report the resulting network. It contains 22 subsectors and 66 edges, with density 0.2857. The mean edge dependency is 0.0566 and the median edge dependency is 0.0561. These values are lower than typical raw asset-level correlations because the network summarizes filtered and aggregated dependencies rather than direct pairwise asset correlations.

The top-degree subsector is Retail, indicating broad direct connectivity in the aggregated dependency map. Electric Utility has the highest betweenness, suggesting that it occupies a bridge-like position in the subsector-level network. Apparel and Footwear has the highest weighted degree, reflecting stronger dependency intensity across its retained connections. These summaries help translate the ticker-level PMFG structure into economic language.

Table 12 gives the node-level details behind this interpretation. Degree measures how many of the 66 retained subsector links touch a given node, whereas weighted degree also accounts for the strength of those links. The two measures therefore need not identify the same subsector: a node may have many relatively weak connections or fewer but stronger ones. Betweenness identifies nodes that connect otherwise less-related parts of the graph, while internal mean dependency summarizes the average filtered dependency among assets belonging to the same subsector. For subsectors represented by a single asset, this within-subsector quantity is reported as zero because no within-subsector pair exists.

Figure 15 provides the complementary structural view. Each node represents a subsector, its colour indicates the macro-sector, and its size is proportional to the number of assets in that subsector. The spatial layout is only a visualization of graph connectivity; distances between plotted nodes should not be interpreted as additional economic distances. Thick edges indicate larger average retained dependencies and make concentrated cross-subsector relationships visible, while bridge-like positions and peripheral nodes can be compared with the betweenness and degree values in Table 12. Because the graph is built after ticker-level filtering and subsector aggregation, it should be read as a compact descriptive summary of the retained dependency structure, not as a causal or point-in-time network.

The subsector network motivates the exploratory retrospective model assessment that follows. It does not establish that centrality predicts future volatility: a real-time test would require the network to be reconstructed at each forecast origin.

Table 12: Subsector-level summary of the aggregated dependency network.

Subsector	Macro-sector	Assets	Degree	Weighted degree	Betweenness	Internal mean dep.
Aerospace and Defense	Industrials	1	4	0.1355	0.0381	0.0000
Apparel and Footwear	Consumer Cyclical	2	12	0.9412	0.0190	0.1676
Banks	Financials	5	5	0.2540	0.0786	0.0841
Chemicals	Basic Materials	3	6	0.2134	0.2833	0.0497
Construction Materials	Industrials	1	9	0.6055	0.0095	0.0000
Electric Utility	Utilities	10	4	0.1080	0.3310	0.1033
Financial Services	Financials	2	4	0.2911	0.0000	0.1096
Gas Utility	Utilities	1	5	0.2563	0.0190	0.0000
Insurance	Financials	1	7	0.3801	0.0190	0.0000
Machinery and Equipment	Industrials	3	8	0.4972	0.0571	0.0853
Mining	Basic Materials	3	4	0.3005	0.0190	0.2330
Oil Gas and Biofuels	Oil Gas and Biofuels	2	6	0.2602	0.1214	0.0714
Paper and Cellulose	Basic Materials	1	4	0.2082	0.0571	0.0000
Personal and Home Care	Consumer Non-Cyclical	1	5	0.2640	0.0048	0.0000
Real Estate	Real Estate	2	4	0.1561	0.0476	0.0093
Retail	Consumer Cyclical	2	13	0.8550	0.0667	0.1202
Steel and Metallurgy	Basic Materials	9	4	0.2513	0.0857	0.1161
Telecommunications	Telecommunications	2	4	0.1204	0.0714	0.7955
Textiles	Consumer Cyclical	1	7	0.4747	0.0476	0.0000
Transportation	Industrials	2	8	0.3689	0.3214	0.0428
Transportation Equipment	Industrials	2	5	0.3595	0.0000	0.0584
Water and Sanitation	Utilities	2	4	0.1738	0.0143	0.0691

Notes: The network contains 22 subsectors and 66 retained undirected links (density 0.2857); its mean and median edge dependencies are 0.0566 and 0.0561. Degree counts retained subsector links, weighted degree sums their dependency weights, betweenness is the standard shortest-path

RMT-filtered networks: retrospective volatility assessment

measure on the aggregated graph, and internal mean dependency is calculated from asset pairs within each subsector. The edge weights W_{ab} are defined in Section 4.3.2.

Display convention: The global numerical-precision convention in Section 4.4 applies.




Figure 15: Aggregated subsector dependency network. Nodes represent subsectors, node colour identifies the macro-sector and node size is proportional to the number of assets in the subsector. Edges represent retained subsector-level dependency weights, with thicker edges indicating stronger average dependency.

10. Volatility Forecasting

This section evaluates whether the full-sample structural descriptors are associated with model performance under the stated chronological split. It reports model configurations, held-out metrics, time-series checks and feature-importance diagnostics. Because the structural descriptors are ex-post, every result in this section is interpreted as retrospective evidence rather than as a deployable point-in-time forecasting result.

10.1. Chronological evaluation with ex-post structural descriptors

We conduct a retrospective evaluation of whether RMT and network descriptors are associated with realized-volatility model performance. The experiment uses PETR4, VALE3 and BBDC4 as target assets and compares machine-learning models with simple neural extensions. The data are split chronologically: training covers 2006–2017, validation covers 2018–2020 and testing covers 2021–2025. This ordering keeps model fitting and feature-set selection separated from the held-out period. However, the RMT and network descriptors are calculated from the full synchronized panel and broadcast across dates; Sets B and C consequently use ex-post structural information. Their test-period scores are not evidence of a tradable, real-time forecasting advantage. They are an exploratory added-information comparison that a recursive rolling design must validate.

10.2. Machine-learning and deep-learning models

The machine-learning models use three nested feature sets in an ablation design. Set A contains standard lagged return and volatility features. Set B adds rolling market variables and full-sample RMT descriptors. Set C adds full-sample MST and PMFG descriptors. This nesting allows the experiment to assess whether ex-post network structure is associated with model performance beyond lagged-volatility persistence and broad market/RMT descriptors.

The deep-learning extension uses two simple architectures. The MLP is a tabular neural regressor trained on the same feature sets as the machine-learning models. The CNN-1D uses short rolling windows of return and volatility variables. Both neural models use positive-output regression for realized volatility. These models are included as an incremental extension, not as a claim that the neural architecture has been optimized for this market.

10.3. Chronological test-period model comparison

Figure 16 summarizes held-out test-period QLIKE loss with moving-block bootstrap intervals, and Table 13 reports the aggregated test-period metrics. Lower QLIKE is better. For each model and horizon, the reported feature set is selected using the validation period and then evaluated on the test period. The temporal holdout applies to model fitting and selection; it does not make the full-sample structural descriptors point-in-time features.




Figure 16: Chronological test-period comparison for machine-learning and deep-learning volatility models. Points are mean QLIKE across PETR4, VALE3 and BBDC4; horizontal bars are 95% circular moving-block bootstrap intervals with block length equal to the horizon. Labels show the validation-selected feature set for each model and horizon. Sets B and C include ex-post full-sample structural descriptors and are not real-time forecasting specifications.

At the 5-day horizon, Random Forest with Set C obtains the lowest average test-period QLIKE among the validation-selected models. Ridge, CNN-1D and MLP remain close in MAE and RMSE, and the 95% moving-block QLIKE intervals overlap materially across the leading specifications. Thus the ranking is descriptive rather than evidence of a precisely estimated performance separation. Per-asset metrics and intervals are reported in the supplementary material.

At the 20-day horizon, Ridge with Set C obtains the lowest average test-period QLIKE, while CNN-1D gives competitive MAE, RMSE and held-out R^2 . The MLP has relatively low MAE and RMSE but weaker QLIKE, indicating that it can track average volatility levels while still producing less favourable variance-proxy loss. The focused HAC loss-difference check in Table 14 compares the validation-selected structural champion (Ridge Set C at both horizons) with the validation-selected point-in-time Set A baseline (Ridge). The mean baseline-minus-structural difference is negative at both horizons, and the one-sided tests do not support lower structural QLIKE. Accordingly, the retrospective exercise does not establish a forecasting gain from the ex-post descriptors; a recursive feature construction remains necessary for that claim.

Table 13: Aggregated chronological test-period comparison (2021–2025). Values are averaged across PETR4, VALE3 and BBDC4 after validation-based model selection. Sets B and C use ex-post full-sample structural descriptors.

Horizon	Model	Feature set	MAE	RMSE	QLIKE	95% CI	R^2_{oos}
5	Random Forest	Set C (+Network)	0.0161	0.0227	-5.2508	[-5.3859, -5.0970]	0.1989
5	Ridge	Set C (+Network)	0.0146	0.0214	-5.2303	[-5.3787, -5.0652]	0.2774
5	HistGradientBoosting	Set C (+Network)	0.0168	0.0234	-5.2228	[-5.3580, -5.0691]	0.1477
5	CNN-1D	Sequence window	0.0143	0.0213	-5.2098	[-5.3634, -5.0367]	0.2886
5	MLP	Set A (Classical)	0.0138	0.0210	-5.2095	[-5.3658, -5.0292]	0.3046
20	Ridge	Set C (+Network)	0.0219	0.0309	-3.8812	[-4.0443, -3.7206]	0.3803
20	CNN-1D	Sequence window	0.0205	0.0306	-3.8552	[-4.0470, -3.6445]	0.3951
20	Random Forest	Set C (+Network)	0.0264	0.0368	-3.8455	[-3.9911, -3.6843]	0.1667
20	HistGradientBoosting	Set C (+Network)	0.0265	0.0371	-3.8417	[-3.9841, -3.6847]	0.1685
20	MLP	Set B (+Market/RMT)	0.0213	0.0305	-3.4271	[-3.8434, -2.8528]	0.3726

Notes: MAE is mean absolute error, RMSE is root mean squared error, QLIKE is the variance-proxy loss defined in Section 4.4, and R^2_{oos} is computed relative to the training-sample mean benchmark. The QLIKE interval is a 95% circular moving-block bootstrap interval based on 2,000 replications and a block length equal to the horizon. Set A contains point-in-time return and volatility predictors. Set B adds rolling market variables and full-sample RMT descriptors; Set C adds full-sample MST/PMFG descriptors. Thus the table evaluates chronological model fitting and selection, not a strict real-time forecast implementation for Sets B and C.

Display convention: The global numerical-precision convention in Section 4.4 applies.

Table 14: HAC loss-difference test for the validation-selected structural specification against the validation-selected Set A baseline. Positive differences favour the structural specification.

Horizon	Structural	Set A baseline	Mean difference	HAC s.e.	Statistic	One-sided p
5	Ridge Set C	Ridge Set A	-0.0045	0.0031	-1.47	0.929
20	Ridge Set C	Ridge Set A	-0.0072	0.0039	-1.86	0.969

Notes: The loss differential is the date-level mean of $QLIKE_t^A - QLIKE_t^{\text{struct}}$ across the three assets. The HAC variance uses Bartlett/Newey–West lags 4 and 19 for the 5- and 20-day horizons, respectively. The one-sided alternative is that the structural loss is lower.

Display convention: The global numerical-precision convention in Section 4.4 applies.

10.4. Realized versus predicted volatility checks

Aggregate metrics can hide important time-series behaviour. Figure 17 compares realized 20-day volatility with the validation-selected machine-learning or deep-learning forecast for PETR4, VALE3 and BBDC4. We include this plot in the main text because it illustrates how the selected model tracks changing volatility regimes.

The plots suggest that broad volatility regimes are tracked reasonably well, but extreme spikes remain difficult to forecast. This matches the aggregate results: flexible models can be competitive, but lagged volatility remains the main source of point-in-time predictive information. The network variables should be read as ex-post structural descriptors rather than replacements for time-series volatility information.




Figure 17: Realized versus predicted 20-day volatility for the validation-selected machine-learning or deep-learning model. The vertical dashed line marks the beginning of the 2021–2025 test period.

10.5. Feature importance and network signal

Random Forest feature importance gives a model-specific check of which predictors contribute most to the fitted model. Figure 18 shows that standard volatility features dominate the model. Rolling absolute returns, lagged realized volatility, rolling volatility and EWMA variables carry most of the importance.

The strongest features are rolling absolute return over 20 days, lagged realized volatility and longer rolling volatility measures. Network descriptors have smaller importance values than standard volatility predictors, but selected PMFG variables enter the top-ranked predictors, especially at the 20-day horizon. Because Random Forest importance measures can be affected by correlation among predictors, and because the structural descriptors are ex post, these rankings are model checks rather than causal or real-time measures of variable relevance. In summary, topology does not replace lagged-volatility features; its point-in-time contribution remains to be tested with recursively reconstructed features.




Figure 18: Feature importance for Random Forest volatility forecasting models. Standard volatility predictors dominate the selected rankings, while selected PMFG-based and sector/subsector variables appear with smaller but nonzero importance in the network-augmented feature set.

11. Discussion

Taken together, the results point to a layered interpretation of Brazilian equity dependencies, in line with the multi-scale perspective advocated by Raddant and Di Matteo [42]. The first layer is a market-mode component. It is reflected in the moderate positive average correlation ($\bar{\rho} \approx 0.24$), the increase in rolling average correlation around stress periods and the very large leading eigenvalue ($\lambda_1 = 21.65$, approximately 37.3% of the trace of the correlation matrix). We interpret this component as broad market synchronization. The analysis does not identify the economic mechanisms behind that synchronization; its role here is to show why raw networks can obscure weaker localized patterns.

The second layer consists of candidate group and sector structure. Within-sector correlations are higher than between-sector correlations on average, and the RMT decomposition identifies four additional eigenvalues above the random upper edge. Eigenvector loadings and group-mode matrices suggest that commodity, financial, telecommunications and utility-related structures contribute to the non-random part of the spectrum. This reading is descriptive, not causal: we identify structure in the dependency matrix, not the economic mechanisms that generate each edge.

The third layer is noise-compatible residual dependence. Not every correlation is economically meaningful, and not every edge retained by a graph is a stable economic relationship. This is why we combine RMT, heatmaps, dendrograms and network metrics. A result is more credible when it appears across several representations, not only in one visualization.

The network results show that market-mode removal changes structure in a systematic way. Original networks have stronger average edge correlations, while group-mode networks have much lower average edge correlations. However, the group-mode PMFG has higher clustering (0.6653 versus 0.5273). Notably, lower average edge strength does not mean absence of structure. It means that the dominant common component has been removed and localized organization becomes easier to observe [15, 16]. The PMFG is well suited for this analysis because it preserves richer graph structure than the MST, retaining 168 edges with 166 triangles and 55 4-cliques that allow hub-rank shifts, clustering changes and subsector dependency maps to be studied in detail.

Clustering summarizes hierarchical organization, but it does not identify graph-theoretic hubs, bridges or local cycles. The transition from dendrograms to filtered financial networks adds information. Hub reorganization after RMT filtering is meaningful: original hubs such as GGBR4 and BBDC4 are central in networks that include the common component, while group-mode hubs such as GUAR3 occupy central positions after the leading component is removed. This suggests that centrality in financial networks depends on the matrix used to construct the network.

The volatility-modelling results require caution. Machine-learning and simple deep-learning models do not produce uniform differences across horizons and metrics. Instead, models augmented with ex-post network descriptors are competitive at selected horizons, while lagged volatility variables remain dominant. Because the RMT and network descriptors are estimated from the full sample, these results identify a retrospective structural association rather than

a real-time forecast gain. The net interpretation is modest: topology summarizes cross-sectional geometry that may complement the temporal dynamics of individual asset variance, but its point-in-time value remains untested.

The broader contribution is an integrated and reproducible framework connecting econophysics filtering, financial networks and volatility modelling for B3 equities. Separating chronological model fitting from ex-post structural enrichment makes the current evidence easier to interpret and defines the next step: reconstruct RMT and network quantities in rolling or expanding windows at every forecast origin.

12. Limitations

The results should be read subject to five limits.

1. **Information availability.** RMT and network descriptors are estimated from the full synchronized sample and are therefore unavailable at historical forecast origins. The volatility exercise is retrospective, not a strict real-time test.
2. **Target coverage.** The model assessment uses PETR4, VALE3 and BBDC4. These liquid and economically distinct assets are useful demonstrations, but do not represent the whole B3 cross-section.
3. **Historical data quality.** Long adjusted-price histories can retain artifacts from corporate actions, ticker transitions, units or adjustment factors. Broader universes require renewed event-level validation.
4. **Classification and issuer representation.** Sector and subsector labels are an initial descriptive map, and the deterministic issuer-collapse rule is a robustness check rather than a definitive economic-exposure or liquidity rule.
5. **Model scope.** Static network descriptors and simple neural architectures do not establish causal mechanisms. Future work should reconstruct structural features in rolling or expanding windows, broaden the target set and compare richer temporal or graph-based models.

13. Conclusion

We examined the dependency structure of Brazilian equities traded on B3 using an econophysics workflow. The raw historical universe covers 1998–2025; the RMT, clustering and network findings are based on its complete 58-asset panel from 23 August 2006 to 15 August 2024. The main findings can be summarized along four dimensions.

First, the stylized-fact and correlation analysis shows that Brazilian equities display signatures commonly associated with complex financial systems: heavy-tailed return behaviour, volatility clustering, persistent absolute-return dependence and a moderate but positive average cross-correlation ($\bar{\rho} \approx 0.24$). Within-sector correlations are larger than between-sector correlations on average, confirming that sectoral organization is reflected in the dependency matrix.

Second, Random Matrix Theory decomposes the correlation matrix. Five eigenvalues exceed the Marčenko–Pastur upper bound, with the largest eigenvalue accounting for approximately 37.3% of the trace of the correlation matrix. The leading component captures a market-mode component, while the remaining eigenvectors outside the random band suggest commodity, financial and utility-related group structures. The reconstruction checks indicate numerical accuracy within floating-point precision, and the group-mode matrix retains localized dependency patterns after the dominant market component is removed.

Third, financial network analysis through MST and PMFG representations shows that RMT filtering changes the structure of asset dependencies. The group-mode PMFG has lower mean edge correlation but higher clustering than the original PMFG, indicating stronger local clustering after market-mode removal. Hub rankings shift from networks dominated by broad market co-movement to networks where different assets occupy central positions after excluding the leading component. The PMFG retains 168 edges with triangles and 4-cliques, leaving local graph structures that are absent from the MST and useful for subsector-level analysis.

Fourth, the chronological volatility-model assessment identifies a limited ex-post association between network topology and model performance. Some model–horizon comparisons favor full-sample structural descriptors, while the CNN-1D model is competitive in MAE, RMSE and held-out R^2 . Lagged-volatility features remain dominant, and selected PMFG centralities have smaller importance. These comparisons do not establish a real-time forecasting benefit because the RMT and network descriptors use full-sample information.

The contribution is therefore structural and methodological: network descriptors summarize the cross-sectional geometry of the market but do not replace lagged-volatility dynamics. The essential next step is to construct RMT

and network features recursively at each forecast origin, then repeat model selection and evaluation without future information. That design can determine whether the retrospective structural association survives as a real-time forecasting benefit.

A. Supplementary figures

Additional figures and tables generated by the computational workflow are retained as supplementary material. These include asset-universe tables, return summaries, correlation matrices, RMT eigenvector loadings, network edge lists, clique lists, centrality rankings and detailed forecasting results by asset, horizon, model and feature set.

Page 35 of 42

RMT-filtered networks: retrospective volatility assessment




Figure 20: Full RMT decomposition: original correlation, market-mode, candidate group/sector, noise-compatible and retained components. The main text presents the original and group/sector panels at readable scale.

B. Supplementary tables

The following tables are generated from the computational outputs by the supplementary-table generation script. They provide the complete classified asset universe, the deterministic issuer/share-class selection, spectral loadings, null tests and asset-level forecast metrics used in the manuscript.

All numerical entries below are generated from the CSV files named in the captions.

B.1. Asset universe

Table 15: Asset universe

Symbol	Company	Macro-sector	Subsector
ALPA4	ALPARGATAS	Consumer Cyclical	Apparel and Footwear
BAZA3	BANCO DA AMAZONIA	Financials	Banks
BBAS3	BANCO DO BRASIL	Financials	Banks
BBDC3	BRADESCO	Financials	Banks
BBDC4	BRADESCO	Financials	Banks
BEES3	BANESTES	Financials	Banks
BOBR4	BOMBRIL	Consumer Non-Cyclical	Personal and Home Care
BRAP3	BRADESPAR	Basic Materials	Mining
BRAP4	BRADESPAR	Basic Materials	Mining
BRKM3	BRASKEM	Basic Materials	Chemicals
BRKM5	BRASKEM	Basic Materials	Chemicals
CCRO3	CCR	Industrials	Transportation
CGAS5	COMGAS	Utilities	Gas Utility
CMIG3	CEMIG	Utilities	Electric Utility
CMIG4	CEMIG	Utilities	Electric Utility
COCE5	COELCE	Utilities	Electric Utility
CPFE3	CPFL ENERGIA	Utilities	Electric Utility
CPLE3	COPEL	Utilities	Electric Utility
CPLE6	COPEL	Utilities	Electric Utility
CSNA3	CSN	Basic Materials	Steel and Metallurgy
CTNM4	COTEMINAS	Consumer Cyclical	Textiles
CYRE3	CYRELA	Real Estate	Real Estate
ELET3	ELETROBRAS	Utilities	Electric Utility
ELET6	ELETROBRAS	Utilities	Electric Utility
EMAE4	EMAE	Utilities	Electric Utility
EMBR3	EMBRAER	Industrials	Aerospace and Defense
ETER3	ETERNIT	Industrials	Construction Materials
FESA4	FERBASA	Basic Materials	Steel and Metallurgy
GGBR3	GERDAU	Basic Materials	Steel and Metallurgy
GGBR4	GERDAU	Basic Materials	Steel and Metallurgy
GOAU3	METALURGICA GERDAU	Basic Materials	Steel and Metallurgy
GOAU4	METALURGICA GERDAU	Basic Materials	Steel and Metallurgy
GRND3	GRENDENE	Consumer Cyclical	Apparel and Footwear

Symbol	Company	Macro-sector	Subsector
GUAR3	GUARARAPES	Consumer Cyclical	Retail
ITSA3	ITAUSA	Financials	Financial Services
ITSA4	ITAUSA	Financials	Financial Services
KEPL3	KEPLER WEBER	Industrials	Machinery and Equipment
KLBN4	KLABIN	Basic Materials	Paper and Cellulose
LREN3	LOJAS RENNER	Consumer Cyclical	Retail
MGEL4	MANGELS	Basic Materials	Steel and Metallurgy
PETR3	PETROBRAS	Oil Gas and Biofuels	Oil Gas and Biofuels
PETR4	PETROBRAS	Oil Gas and Biofuels	Oil Gas and Biofuels
POMO4	MARCOPOLO	Industrials	Transportation Equipment
PSSA3	PORTO SEGURO	Financials	Insurance
RAPT4	RANDON	Industrials	Transportation Equipment
RENT3	LOCALIZA	Industrials	Transportation
ROMI3	INDS ROMI	Industrials	Machinery and Equipment
SAPR4	SANEPAR	Utilities	Water and Sanitation
SBSP3	SABESP	Utilities	Water and Sanitation
SCAR3	SAO CARLOS	Real Estate	Real Estate
SHUL4	SCHULZ	Industrials	Machinery and Equipment
TELB3	TELEBRAS	Telecommunications	Telecommunications
TELB4	TELEBRAS	Telecommunications	Telecommunications
TRPL4	ISA CTEEP	Utilities	Electric Utility
UNIP6	UNIPAR	Basic Materials	Chemicals
USIM3	USIMINAS	Basic Materials	Steel and Metallurgy
USIM5	USIMINAS	Basic Materials	Steel and Metallurgy
VALE3	VALE	Basic Materials	Mining

B.2. Issuer/share-class selection

Table 16: Issuer/share-class selection

Company	Symbol	Valid returns	Retained
BRADESCO	BBDC3	6885	No
BRADESCO	BBDC4	6886	Yes
BRADESPAR	BRAP3	6210	No
BRADESPAR	BRAP4	6291	Yes
BRASKEM	BRKM3	5048	No
BRASKEM	BRKM5	5783	Yes
CEMIG	CMIG3	6880	No
CEMIG	CMIG4	6885	Yes
COPEL	CPLE3	6787	No
COPEL	CPLE6	6852	Yes
ELETROBRAS	ELET3	6852	Yes
ELETROBRAS	ELET6	6852	No
GERDAU	GGBR3	5682	No
GERDAU	GGBR4	6174	Yes
ITAUSA	ITSA3	5710	No
ITAUSA	ITSA4	6885	Yes
METALURGICA GERDAU	GOAU3	5852	No
METALURGICA GERDAU	GOAU4	6885	Yes
PETROBRAS	PETR3	6885	Yes
PETROBRAS	PETR4	6885	No
TELEBRAS	TELB3	5386	No
TELEBRAS	TELB4	6852	Yes
USIMINAS	USIM3	6005	No

Company	Symbol	Valid returns	Retained
USIMINAS	USIM5	6669	Yes

B.3. Largest absolute eigenvector loadings

Table 17: Largest absolute eigenvector loadings

Rank	Eigenvalue	Symbol	Sector	$ u $
1	21.6505	GOAU4	Basic Materials	0.1743
1	21.6505	GGBR4	Basic Materials	0.1713
1	21.6505	ITSA4	Financials	0.1694
1	21.6505	BBDC3	Financials	0.1676
1	21.6505	BBDC4	Financials	0.1675
2	3.0957	VALE3	Basic Materials	0.2831
2	3.0957	BRAP4	Basic Materials	0.2831
2	3.0957	BRAP3	Basic Materials	0.2507
2	3.0957	GGBR4	Basic Materials	0.2487
2	3.0957	GGBR3	Basic Materials	0.2387
3	1.7839	ROMI3	Industrials	0.2922
3	1.7839	SHUL4	Industrials	0.2903
3	1.7839	BBDC4	Financials	0.2735
3	1.7839	BBDC3	Financials	0.2689
3	1.7839	ITSA4	Financials	0.2623
4	1.6543	TELB3	Telecommunications	0.6936
4	1.6543	TELB4	Telecommunications	0.6882
4	1.6543	GRND3	Consumer Cyclical	0.0859
4	1.6543	ITSA3	Financials	0.0808
4	1.6543	MGEL4	Basic Materials	0.0694
5	1.4835	CMIG3	Utilities	0.2901
5	1.4835	CMIG4	Utilities	0.2577
5	1.4835	CPLE6	Utilities	0.2502
5	1.4835	CPLE3	Utilities	0.2265
5	1.4835	ALPA4	Consumer Cyclical	0.2129

B.4. Circular-shift RMT null test

Table 18: Circular-shift RMT null test

Rank	Empirical λ	Null 95th pct.	Exceedances	p
1	21.6505	1.6882	0	0.0010
2	3.0957	1.4932	0	0.0010
3	1.7839	1.3889	0	0.0010
4	1.6543	1.3533	0	0.0010
5	1.4835	1.3295	0	0.0010

B.5. Asset-level sector-label permutation test

Table 19: Asset-level sector-label permutation test

Assets	Pairs	Permutations	Δ_{sector}	Null 95th pct.	p
58	1653	10000	0.1184	0.0245	0.0001

B.6. Forecast metrics by asset and horizon

Table 20: Forecast metrics by asset and horizon

Horizon	Model	Set	Asset	MAE	RMSE	QLIKE [95% CI]
5	CNN-1D	Seq.	BBDC4	0.0140	0.0226	-5.1902 [-5.4841, -4.8206]
5	CNN-1D	Seq.	PETR4	0.0156	0.0240	-5.0598 [-5.2875, -4.7550]
5	CNN-1D	Seq.	VALE3	0.0134	0.0172	-5.3793 [-5.5121, -5.2431]
5	HistGradientBoosting	Set C C	BBDC4	0.0160	0.0239	-5.2507 [-5.4862, -4.9763]
5	HistGradientBoosting	Set C C	PETR4	0.0184	0.0264	-5.0532 [-5.2719, -4.7793]
5	HistGradientBoosting	Set C C	VALE3	0.0159	0.0199	-5.3646 [-5.4702, -5.2573]
5	MLP	Set A A	BBDC4	0.0135	0.0226	-5.1709 [-5.4817, -4.7816]
5	MLP	Set A A	PETR4	0.0154	0.0238	-5.0529 [-5.2731, -4.7620]
5	MLP	Set A A	VALE3	0.0126	0.0166	-5.4045 [-5.5466, -5.2630]
5	Random Forest	Set C C	BBDC4	0.0157	0.0234	-5.2852 [-5.5174, -5.0005]
5	Random Forest	Set C C	PETR4	0.0176	0.0259	-5.0806 [-5.2863, -4.7902]
5	Random Forest	Set C C	VALE3	0.0151	0.0187	-5.3866 [-5.4934, -5.2723]
5	Ridge	Set C C	BBDC4	0.0148	0.0232	-5.1685 [-5.4570, -4.8014]
5	Ridge	Set C C	PETR4	0.0157	0.0239	-5.0905 [-5.3091, -4.8093]
5	Ridge	Set C C	VALE3	0.0134	0.0170	-5.4319 [-5.5517, -5.3145]
20	CNN-1D	Seq.	BBDC4	0.0209	0.0325	-3.8624 [-4.1466, -3.4803]
20	CNN-1D	Seq.	PETR4	0.0228	0.0366	-3.6812 [-3.9561, -3.3330]
20	CNN-1D	Seq.	VALE3	0.0179	0.0227	-4.0220 [-4.1925, -3.8565]
20	HistGradientBoosting	Set C C	BBDC4	0.0237	0.0334	-3.9196 [-4.1527, -3.6527]
20	HistGradientBoosting	Set C C	PETR4	0.0296	0.0426	-3.6763 [-3.8730, -3.4413]
20	HistGradientBoosting	Set C C	VALE3	0.0262	0.0354	-3.9290 [-4.0547, -3.7984]
20	MLP	Set B B	BBDC4	0.0252	0.0355	-2.4923 [-3.6616, -0.6852]
20	MLP	Set B B	PETR4	0.0217	0.0338	-3.7469 [-3.9682, -3.4877]
20	MLP	Set B B	VALE3	0.0171	0.0223	-4.0420 [-4.2032, -3.8846]
20	Random Forest	Set C C	BBDC4	0.0262	0.0358	-3.8826 [-4.1009, -3.6177]
20	Random Forest	Set C C	PETR4	0.0275	0.0407	-3.7113 [-3.9222, -3.4600]
20	Random Forest	Set C C	VALE3	0.0254	0.0339	-3.9425 [-4.0739, -3.8066]
20	Ridge	Set C C	BBDC4	0.0237	0.0338	-3.8738 [-4.1222, -3.6050]
20	Ridge	Set C C	PETR4	0.0227	0.0350	-3.7405 [-3.9609, -3.4636]
20	Ridge	Set C C	VALE3	0.0194	0.0239	-4.0293 [-4.1752, -3.8752]

Funding

This research did not receive any specific grant from funding agencies in the public, commercial, or not-for-profit sectors.

Declaration of competing interest

The authors declare that they have no known competing financial interests or personal relationships that could have appeared to influence the work reported in this paper.

Data availability

The workflow uses locally processed B3/BovDBv2 daily adjusted-price records. The public BovDBv2 dataset is available through Harvard Dataverse under DOI 10.7910/DVN/TMB4IG. The analysis scripts, configuration files and reproducibility documentation are available in the public research-software repository [19]. Raw market-data redistribution remains subject to the terms of the original data source.

Declaration of generative AI and AI-assisted technologies in the manuscript preparation process

During the preparation of this work, the authors used ChatGPT for language refinement, structural organization and code-review support. After using this tool, the authors reviewed and edited the content as needed and take full responsibility for the content of the published article.

CRedit authorship contribution statement

Andrey V. S. Souza: Conceptualization, Methodology, Software, Formal analysis, Writing - original draft. **Bruno L. Dalmazo:** Methodology, Supervision, Writing - review and editing. **Giancarlo Lucca:** Methodology, Validation, Writing - review and editing. **Eduardo N. Borges:** Supervision, Methodology, Writing - review and editing. **Richard F. Pinto:** Software, Data curation, Validation. **Fabian C. Cardoso:** Data curation, Resources, Writing - review and editing. **Viviane L. D. de Mattos:** Formal analysis, Supervision, Writing - review and editing. **Rafael A. Berri:** Supervision, Project administration, Writing - review and editing.

References

- [1] R. N. Mantegna, H. E. Stanley, Introduction to Econophysics: Correlations and Complexity in Finance, Cambridge University Press, Cambridge, 1999.
- [2] R. Cont, Empirical properties of asset returns: stylized facts and statistical issues, Quantitative Finance 1 (2) (2001) 223–236. doi:10.1080/713665670.
- [3] L. Laloux, P. Cizeau, J.-P. Bouchaud, M. Potters, Noise dressing of financial correlation matrices, Physical Review Letters 83 (7) (1999) 1467–1470. doi:10.1103/PhysRevLett.83.1467.
- [4] V. Plerou, P. Gopikrishnan, B. Rosenow, L. A. N. Amaral, T. Guhr, H. E. Stanley, Random matrix approach to cross correlations in financial data, Physical Review E 65 (6) (2002) 066126. doi:10.1103/PhysRevE.65.066126.
- [5] J. Bun, J.-P. Bouchaud, M. Potters, Cleaning large correlation matrices: Tools from random matrix theory, Physics Reports 666 (2017) 1–109. doi:10.1016/j.physrep.2016.10.005.
- [6] M. Billio, M. Getmansky, A. W. Lo, L. Pelizzon, Econometric measures of connectedness and systemic risk in the finance and insurance sectors, Journal of Financial Economics 104 (3) (2012) 535–559. doi:10.1016/j.jfineco.2011.12.010.
- [7] L. S. Junior, I. T. P. Franca, Correlation of financial markets in times of crisis, Physica A: Statistical Mechanics and its Applications 391 (1–2) (2012) 187–208. doi:10.1016/j.physa.2011.07.023.
- [8] Z. Zheng, B. Podobnik, L. Feng, B. Li, Changes in cross-correlations as an indicator for systemic risk, Scientific Reports 2 (2012) 888. doi:10.1038/srep00888.
- [9] B. M. Tabak, D. O. Cajueiro, M. V. Serra, Topology of the brazilian interbank network, Physica A: Statistical Mechanics and its Applications 389 (22) (2010) 5277–5283. doi:10.1016/j.physa.2010.07.037.
- [10] E. J. A. L. Pereira, M. F. da Silva, H. B. B. Pereira, Multiscale network for twenty stock markets using DCCA, Physica A: Statistical Mechanics and its Applications 529 (2019) 121542. doi:10.1016/j.physa.2019.121542.
- [11] T. Stosić, L. Telesca, T. A. E. Ferreira, B. D. Stosic, Correlations and cross-correlations in the Brazilian agrarian commodities and stocks, Physica A: Statistical Mechanics and its Applications 437 (2015) 248–262. doi:10.1016/j.physa.2015.05.091.
- [12] M. A. Leonel Caetano, T. Yoneyama, Complex network analysis of the Brazilian stock market, Expert Systems with Applications 196 (2022) 116680. doi:10.1016/j.eswa.2022.116680.
- [13] V. A. Marčenko, L. A. Pastur, Distribution of eigenvalues for some sets of random matrices, Mathematics of the USSR-Sbornik 1 (4) (1967) 457–483. doi:10.1070/SM1967v001n04ABEH001994.
- [14] R. N. Mantegna, Hierarchical structure in financial markets, The European Physical Journal B 11 (1) (1999) 193–197. doi:10.1007/s100510050929.
- [15] M. Tumminello, T. Aste, T. Di Matteo, R. N. Mantegna, A tool for filtering information in complex systems, Proceedings of the National Academy of Sciences 102 (30) (2005) 10421–10426. doi:10.1073/pnas.0500298102.
- [16] M. Tumminello, F. Lillo, R. N. Mantegna, Correlation, hierarchies, and networks in financial markets, Journal of Economic Behavior & Organization 75 (1) (2010) 40–58. doi:10.1016/j.jebo.2010.01.004.
- [17] F. Pozzi, T. Di Matteo, T. Aste, Spread of risk across financial markets: better to invest in the peripheries, Scientific Reports 3 (2013) 1665. doi:10.1038/srep01665.
- [18] T. Aste, W. Shaw, T. Di Matteo, Correlation structure and dynamics in volatile markets, New Journal of Physics 12 (8) (2010) 085009. doi:10.1088/1367-2630/12/8/085009.
- [19] R. P. N. V. F. Contributors, Rmt planar networks & volatility forecasting — brazil, version 0.1.0, computer software (2026). URL https://github.com/dreyvinixz/RMT_Planar_Networks_Volatility_Forecasting_Brazil
- [20] J.-P. Bouchaud, M. Potters, Theory of financial risk and derivative pricing: From statistical physics to risk management, Cambridge University Press.
- [21] J. Y. Campbell, A. W. Lo, A. C. MacKinlay, The Econometrics of Financial Markets, Princeton University Press, Princeton, NJ, 1997.
- [22] R. S. Tsay, Analysis of Financial Time Series, 2nd Edition, John Wiley & Sons, Hoboken, NJ, 2005.
- [23] B. Mandelbrot, The variation of certain speculative prices, The Journal of Business 36 (4) (1963) 394–419. doi:10.1086/294632.

- [24] E. F. Fama, The behavior of stock-market prices, *The Journal of Business* 38 (1) (1965) 34–105. doi:10.1086/294743.
- [25] R. F. Engle, Autoregressive conditional heteroscedasticity with estimates of the variance of United Kingdom inflation, *Econometrica* 50 (4) (1982) 987–1007. doi:10.2307/1912773.
- [26] T. Bollerslev, Generalized autoregressive conditional heteroskedasticity, *Journal of Econometrics* 31 (3) (1986) 307–327. doi:10.1016/0304-4076(86)90063-1.
- [27] F. Longin, B. Solnik, Extreme correlation of international equity markets, *The Journal of Finance* 56 (2) (2001) 649–676. doi:10.1111/0022-1082.00340.
- [28] A. Ang, J. Chen, Asymmetric correlations of equity portfolios, *Journal of Financial Economics* 63 (3) (2002) 443–494. doi:10.1016/S0304-405X(02)00068-5.
- [29] T. Guhr, A. Müller-Groeling, H. A. Weidenmüller, Random-matrix theories in quantum physics: common concepts, *Physics Reports* 299 (4–6) (1998) 189–425. doi:10.1016/S0370-1573(97)00088-4.
- [30] G. Bonanno, G. Caldarelli, F. Lillo, S. Micciché, N. Vandewalle, R. N. Mantegna, Networks of equities in financial markets, *The European Physical Journal B* 38 (2) (2004) 363–371. doi:10.1140/epjb/e2004-00129-6.
- [31] T. G. Andersen, T. Bollerslev, F. X. Diebold, P. Labys, Modeling and forecasting realized volatility, *Econometrica* 71 (2) (2003) 579–625. doi:10.1111/1468-0262.00418.
- [32] A. E. Hoerl, R. W. Kennard, Ridge regression: Biased estimation for nonorthogonal problems, *Technometrics* 12 (1) (1970) 55–67. doi:10.1080/00401706.1970.10488634.
- [33] L. Breiman, Random forests, *Machine Learning* 45 (1) (2001) 5–32. doi:10.1023/A:1010933404324.
- [34] J. H. Friedman, Greedy function approximation: A gradient boosting machine, *Annals of Statistics* 29 (5) (2001) 1189–1232. doi:10.1214/aos/1013203451.
- [35] A. Bucci, Realized volatility forecasting with neural networks, *Journal of Financial Econometrics* 18 (3) (2020) 502–531. doi:10.1093/jjfinec/nbaa008.
- [36] K. Christensen, M. Siggard, B. Veliyev, A machine learning approach to volatility forecasting, *Journal of Financial Econometrics* 21 (5) (2023) 1680–1727. doi:10.1093/jjfinec/nbac020.
- [37] A. J. Patton, Volatility forecast comparison using imperfect volatility proxies, *Journal of Econometrics* 160 (1) (2011) 246–256. doi:10.1016/j.jeconom.2010.03.034.
- [38] A. Smolyak, S. Havlin, Three decades in econophysics—from microscopic modelling to macroscopic complexity and back, *Entropy* 24 (2) (2022) 271. doi:10.3390/e24020271.
- [39] M. Vogl, Controversy in financial chaos research and nonlinear dynamics: A short literature review, *Chaos, Solitons & Fractals* 162 (2022) 112444. doi:10.1016/j.chaos.2022.112444.
- [40] K. T. Chi, J. Liu, F. C. M. Lau, A network perspective of the stock market, *Journal of Empirical Finance* 17 (4) (2010) 659–667. doi:10.1016/j.jempfin.2010.04.008.
- [41] S. Kumar, N. Deo, Correlation and network analysis of global financial indices, *Physical Review E* 86 (2) (2012) 026101. doi:10.1103/PhysRevE.86.026101.
- [42] M. Raddant, T. Di Matteo, A look at financial dependencies by means of econophysics and financial economics, *Journal of Economic Interaction and Coordination* 18 (4) (2023) 701–734. doi:10.1007/s11403-023-00395-y.
- [43] L. Laloux, P. Cizeau, M. Potters, J.-P. Bouchaud, Random matrix theory and financial correlations, *International Journal of Theoretical and Applied Finance* 3 (3) (2000) 391–397. doi:10.1142/S0219024900000255.
- [44] J.-P. Onnela, A. Chakraborti, K. Kaski, J. Kertész, A. Kanto, Dynamics of market correlations: Taxonomy and portfolio analysis, *Physical Review E* 68 (5) (2003) 056110. doi:10.1103/PhysRevE.68.056110.
- [45] D.-M. Song, M. Tumminello, W.-X. Zhou, R. N. Mantegna, Evolution of worldwide stock markets, correlation structure, and correlation-based graphs, *Physical Review E* 84 (2) (2011) 026108. doi:10.1103/PhysRevE.84.026108.
- [46] G. P. Massara, T. Di Matteo, T. Aste, Network filtering for big data: Triangulated maximally filtered graph, *Journal of Complex Networks* 5 (2) (2017) 161–178. doi:10.1093/comnet/cnw015.
- [47] L. Zhao, W. Li, X. Cai, Stock market as temporal network, *Physica A: Statistical Mechanics and its Applications* 506 (2018) 1104–1112. doi:10.1016/j.physa.2018.05.039.
- [48] D. Y. Kenett, M. Tumminello, A. Madi, G. Gur-Gershgoren, R. N. Mantegna, E. Ben-Jacob, Dominating clasp of the financial sector revealed by partial correlation analysis of the stock market, *PLoS ONE* 5 (12) (2010) e15032. doi:10.1371/journal.pone.0015032.
- [49] T. C. Silva, S. R. S. de Souza, B. M. Tabak, Structure and dynamics of the global financial network, *Chaos, Solitons & Fractals* 88 (2016) 218–234. doi:10.1016/j.chaos.2016.01.023.
- [50] G.-J. Wang, C. Xie, H. E. Stanley, Correlation structure and evolution of world stock markets: Evidence from Pearson and partial correlation-based networks, *Computational Economics* 51 (3) (2018) 607–635. doi:10.1007/s10614-016-9627-7.
- [51] P. Ferreira, É. Pereira, M. F. Silva, H. B. B. Pereira, Detrended correlation coefficients between oil and stock markets, *Chaos, Solitons & Fractals* 127 (2019) 120–128. doi:10.1016/j.chaos.2019.06.037.
- [52] S.-H. Poon, C. W. J. Granger, Forecasting volatility in financial markets: A review, *Journal of Economic Literature* 41 (2) (2003) 478–539. doi:10.1257/jel.41.2.478.
- [53] F. Corsi, A simple approximate long-memory model of realized volatility, *Journal of Financial Econometrics* 7 (2) (2009) 174–196. doi:10.1093/jjfinec/nbp001.
- [54] P. R. Hansen, A. Lunde, A forecast comparison of volatility models: Does anything beat a GARCH(1,1)?, *Journal of Applied Econometrics* 20 (7) (2005) 873–889. doi:10.1002/jae.800.
- [55] D. Zhang, D. C. Broadstock, Financial network topology and market risk: A machine learning perspective, *International Review of Financial Analysis* 87 (2023) 102604. doi:10.1016/j.irfa.2023.102604.

- [56] F. C. Cardoso, J. A. V. Malska, P. J. Ramiro, G. Lucca, E. N. Borges, V. L. D. d. Mattos, R. A. Berri, BovDB: a data set of stock prices of all companies in B3 from 1995 to 2020, *Journal of Information and Data Management* 13 (1). doi:10.5753/jidm.2022.2345.
- [57] A. S. Souza, G. Lucca, E. N. Borges, F. C. Cardoso, B. L. Dalmazo, R. Berri, Dataset for Intraday Analysis of B3 stock prices, *Harvard Dataverse*, v1 (2024). doi:10.7910/DVN/TMB4IG.
URL <https://doi.org/10.7910/DVN/TMB4IG>
- [58] R. R. Sokal, F. J. Rohlf, The comparison of dendrograms by objective methods, *Taxon* 11 (2) (1962) 33–40. doi:10.2307/1217208.