

Highlights

Random Matrix Filtering and Planar Financial Networks for Volatility Forecasting in the Brazilian Stock Market

Andrey V. S. Souza, Bruno L. Dalmazo, Giancarlo Lucca, Eduardo N. Borges, Richard F. Pinto, Fabian C. Cardoso, Viviane L. D. de Mattos, Rafael A. Berri

- RMT separates market, group and noise components in B3 equities.
- PMFG filtering highlights localized sectoral dependency patterns.
- Within-sector correlations are larger than between-sector correlations on average.
- Network features add modest but nonzero predictive value for volatility forecasting.
- Simple neural models provide an incremental deep-learning extension.

Random Matrix Filtering and Planar Financial Networks for Volatility Forecasting in the Brazilian Stock Market

Andrey V. S. Souza^{a,*}, Bruno L. Dalmazo^a, Giancarlo Lucca^a, Eduardo N. Borges^a, Richard F. Pinto^a, Fabian C. Cardoso^b, Viviane L. D. de Mattos^a and Rafael A. Berri^a

^aFederal University of Rio Grande (FURG), Rio Grande, Rio Grande do Sul, Brazil

^bUniversity of Rio Verde (UniRV), Rio Verde, Goias, Brazil

ARTICLE INFO

Keywords:

Econophysics
Random matrix theory
Financial networks
Planar maximally filtered graph
Volatility forecasting
Machine learning
Brazilian stock market

ABSTRACT

Financial markets are complex systems in which collective fluctuations, sectoral organization and sampling noise coexist in correlation matrices. We develop an econophysics workflow for Brazilian equities traded on B3 from 1998 to 2025 and test whether the resulting dependency structure contributes to volatility forecasting. Using adjusted daily prices, we compute log returns for a historical universe of 58 equities and combine stylized-fact checks, Random Matrix Theory (RMT), hierarchical clustering, Minimum Spanning Trees (MSTs), Planar Maximally Filtered Graphs (PMFGs) and realized-volatility forecasting models based on machine learning and simple deep-learning architectures. Using the Marčenko–Pastur benchmark, the spectrum is decomposed into a dominant market mode, candidate group/sector modes and a residual noise-compatible component; five eigenvalues lie above the Marčenko–Pastur upper bound, and the largest eigenvalue accounts for about 37% of the trace of the correlation matrix. Filtered MST and PMFG representations show that removing the market mode reduces average edge correlations while retaining localized sectoral patterns. The group-mode PMFG has lower mean edge correlation but higher clustering than the original network, suggesting stronger local clustering in the subsector-level representation. Forecasting experiments for PETR4, VALE3 and BBDC4 indicate that network-augmented Random Forest models are competitive at the 5-day horizon, while Ridge regression and a simple temporal CNN are competitive at the 20-day horizon. Overall, RMT-filtered topology helps interpret market structure and can serve as a complementary input for volatility-risk modelling in emerging equity markets.

1. Introduction

Financial markets are often studied as complex systems in which heterogeneous agents, institutional constraints, liquidity cycles, information flows and feedback mechanisms interact across multiple time scales [1, 2]. Even when the analysis is restricted to daily asset returns, market variation is not fully described by the marginal behaviour of each series. It also has a relational component: how assets move together across the cross-section, summarized by their correlation matrix. This matrix reflects a mixture of collective market motion, sectoral organization, idiosyncratic shocks, crisis amplification and sampling noise. A central task is therefore to separate interpretable dependency structure from fluctuations that are compatible with random estimation effects [3].

The distinction matters because financial risk is inherently multivariate. The volatility of an individual asset matters, but portfolio risk, contagion, diversification and systemic vulnerability depend on how assets co-move [4]. During calm periods, correlations may appear moderate and dispersed. During stress periods, correlations often increase, reducing diversification benefits and lowering the effective dimensionality of the market [5, 6].

Emerging equity markets are useful in this context because they combine global risk channels with local institutional, macroeconomic and liquidity conditions. Brazil is a suitable setting: B3 (Brasil, Bolsa, Balcão), the country's stock exchange and financial-market infrastructure provider,¹ includes globally exposed commodity companies, large

*Corresponding author.

✉ andreyvinicius@furg.br (A.V.S. Souza); dalmazo@furg.br (B.L. Dalmazo); giancarlo.lucca@furg.br (G. Lucca); eduardoborges@furg.br (E.N. Borges); richardpinto@furg.br (R.F. Pinto); fabian@unirv.edu.br (F.C. Cardoso); vivianemattos@furg.br (V.L.D.d. Mattos); rafaelberri@furg.br (R.A. Berri)

ORCID(s): 0009-0006-4177-758X (A.V.S. Souza); 0000-0002-6996-7602 (B.L. Dalmazo); 0000-0002-3776-0260 (G. Lucca); 0000-0003-1595-7676 (E.N. Borges); 0009-0007-0176-3383 (R.F. Pinto); 0000-0002-2842-0387 (F.C. Cardoso); 0000-0002-3512-6290 (V.L.D.d. Mattos); 0000-0002-3812-4186 (R.A. Berri)

¹<https://www.b3.com.br/>

financial institutions, regulated utilities, consumer cyclicals, industrial firms, telecommunications firms and real-estate companies. These assets share exposure to domestic monetary policy, exchange-rate dynamics, country risk and global liquidity, but they also respond to distinct sectoral fundamentals. The dependency matrix therefore combines a dominant market component, group or sectoral components and a large residual component. Despite these features, the Brazilian equity market has received limited attention in the complex-systems literature. Previous studies have addressed Brazilian interbank networks [7], global financial system linkages [8], agrarian cross-correlations [9], multiscale networks [10] and equity-market topology [11], but a comprehensive B3 equity study combining Random Matrix Theory filtering, planar financial networks and volatility forecasting remains insufficiently developed.

Correlation matrices are noisy, especially when the number of assets is not negligible relative to the number of observations. Random Matrix Theory (RMT) offers a benchmark for this problem by comparing the eigenspectrum with the spectrum expected from random correlation matrices [12, 13, 14]. Eigenvalues lying inside the Marčenko–Pastur band are compatible with random sampling noise under the benchmark model, whereas eigenvalues outside the band indicate candidate collective modes that deserve further interpretation. Subsequent work has extended RMT-based methods for cleaning large covariance matrices [15]. We use RMT as an intermediate layer for reconstructing broad market, group-level and noise-compatible components before graph construction, not as a mechanical classifier of economic sectors.

Network filtering gives an additional representation of the same dependency system. By transforming correlations into distances, financial networks retain selected dependency links among assets and make it possible to inspect market backbones, local clusters and central nodes. The Minimum Spanning Tree (MST) emphasizes the sparsest connected structure [16], while the Planar Maximally Filtered Graph (PMFG) retains a denser planar representation with richer local connectivity [17, 18]. We use these constructions to ask whether strong dependencies are organized by sectors, whether central nodes change after market-mode removal and whether the intermediate-scale organization of the market contains information beyond average correlation. Filtered networks are also relevant for risk analysis, where peripheral assets can offer better diversification properties than central hubs [19, 20].

We develop an integrated econophysics workflow for Brazilian equities traded on B3. The workflow has three stages. First, we document stylized facts, static correlations, sectoral correlation differences and rolling market synchronization. Second, we apply RMT to decompose the correlation matrix into market, group/sector and noise components, followed by hierarchical clustering and ordered heatmaps. Third, we convert original and filtered matrices into MST and PMFG networks, study hub reorganization and aggregate dependencies at the subsector level. A volatility forecasting experiment then tests whether the extracted features improve machine-learning and simple deep-learning forecasts under a strict temporal validation design.

The paper makes four contributions:

1. We build a reproducible econophysics workflow for B3 equities using adjusted daily prices from 1998 to 2025, including return construction, stylized-fact checks, correlation analysis, sectoral comparison and rolling synchronization.
2. We apply Random Matrix Theory to a 58-asset historical core universe and reconstruct market, group/sector and noise-compatible components of the empirical correlation matrix, yielding a spectral interpretation of Brazilian equity dependencies.
3. We compare MST and PMFG networks built from original and RMT-filtered matrices, with emphasis on hub reorganization, clique structure, clustering, sectoral retention and subsector-level aggregation.
4. We evaluate whether RMT- and network-derived features improve realized-volatility forecasting in machine-learning and simple deep-learning models, using a strict temporal split and nested feature-set design.

The article is organized as follows. Section 2 reviews financial dependencies, Random Matrix Theory, filtered financial networks and volatility forecasting. Section 3 describes the data and asset universe. Section 4 presents the methodology and the mathematical definitions used throughout the paper. Sections 5–10 report the results. Sections 11 and 12 discuss interpretation and limitations, and Section 13 presents the final considerations.

2. Related work

2.1. Financial dependencies and econophysics

Econophysics studies financial markets using concepts and methods from statistical physics, nonlinear dynamics, complex systems and network science. A central observation is that financial returns exhibit stylized facts that are

difficult to reconcile with independent Gaussian models. Heavy tails, volatility clustering, slow decay in absolute-return autocorrelations, intermittent bursts of activity and collective co-movement patterns are recurrent findings across asset classes, countries and sampling frequencies [1, 2]. These regularities motivate methods that focus not only on the marginal distribution of returns, but also on dependency, synchronization and network structure.

Financial dependency analysis lies at the intersection of econophysics and financial economics. In financial economics, dependence is central to portfolio risk, diversification, contagion, factor structure and systemic vulnerability. In econophysics, the same dependence is often represented through correlation matrices, eigenspectra, filtered graphs and hierarchical structures. Raddant and Di Matteo [3] emphasize that correlations, causality measures, volatility spillovers and networks offer additional views of financial systems. We follow this perspective: the goal is not merely to estimate pairwise correlations, but to decompose the correlation matrix, filter its noisy components, represent the remaining structure as networks and evaluate whether the extracted network structure has forecasting relevance.

Cross-correlation dynamics during crises have received attention because market stress often increases synchronization across assets. Zheng et al. [6] showed that changes in cross-correlation structure can serve as early indicators of systemic instability. Billio et al. [4] proposed econometric measures of connectedness and systemic risk in financial sectors. Junior and Franca [5] documented that correlations among global stock indices tend to increase during crises, reducing the effective dimensionality of international markets. These findings motivate the rolling-correlation and crisis-synchronization analysis undertaken in this paper, where market-wide co-movement is treated as an object of study instead of a nuisance.

The relevance of these studies for the present manuscript is twofold. First, they show that financial dependence is dynamic and crisis-sensitive. Second, they indicate that a single unfiltered correlation matrix may conceal different layers of structure: a broad market component, group or sectoral organization and residual noise. This motivates the use of Random Matrix Theory and filtered networks as additional tools for separating and interpreting these layers.

2.2. Random Matrix Theory in financial correlations

Random Matrix Theory entered finance as a way to separate information from noise in large correlation matrices. Laloux et al. [13, 21] showed that a substantial part of the eigenspectrum of financial correlation matrices can resemble the spectrum predicted for random matrices. The relevant benchmark is the Marčenko–Pastur distribution [12], which sets bounds for the eigenvalues of a random correlation matrix estimated from independent standardized time series. Plerou et al. [14] further studied eigenvalues and eigenvectors outside the random band and read the largest eigenvalue as a market-wide mode.

The intuition is that the empirical correlation matrix contains both structured and noisy components. If N assets are observed over T periods, the ratio $Q = T/N$ determines the width of the random-matrix noise band. Eigenvalues inside this band are treated as compatible with sampling noise, whereas eigenvalues above the upper edge are candidates for collective factors. This interpretation is useful in financial markets because broad market co-movement can dominate the first eigenvalue, while weaker sectoral or group modes may appear in subsequent eigenvectors.

Subsequent work extended RMT-based methods for covariance and correlation cleaning. Bun et al. [15] reviewed and unified several approaches for cleaning large covariance matrices using random matrix theory, connecting theoretical results with practical financial applications. Guhr et al. [22] give a broader mathematical review of random matrix theory and its applications in complex systems. In finance, RMT has been used not only for noise filtering, but also for portfolio optimization, factor identification and the study of collective market behaviour.

Our RMT interpretation is cautious. The largest eigenvalue captures a market mode because it reflects broad positive co-movement across many assets. Eigenvalues above the Marčenko–Pastur upper edge, excluding the first one, are read as group or sector candidates. Eigenvalues inside the random band are treated as noisy residual components. These labels should be read as filtering categories; they do not imply that each eigenvector has a unique economic interpretation. We use this spectral separation as the input for network construction and volatility forecasting in the Brazilian equity market.

2.3. Financial networks: MST, PMFG and filtered graphs

Financial networks transform a dependency matrix into a graph representation. The usual approach begins by converting correlations into distances, so that strongly correlated assets are close in the induced metric space. Mantegna [16] introduced the Minimum Spanning Tree (MST) as a way to reveal hierarchical organization in financial markets. The MST retains the connected tree with minimum total distance and therefore contains exactly $N - 1$ edges for

N assets. Its main advantage is interpretability: it retains a sparse backbone of the strongest dependencies. Its main limitation is that it removes cycles and many potentially meaningful local connections.

Subsequent studies applied MST-based methods to study market taxonomy, instability and dynamics. Onnela et al. [23] analysed the evolution of MST-based market correlations and their implications for portfolio structure. Bonanno et al. [24] extended the network approach to multiple world markets, showing that MSTs often organize assets according to sectoral and geographical structure. Song et al. [25] studied the evolution of worldwide stock markets through correlation-based graphs, documenting time variation in network structure.

The Planar Maximally Filtered Graph (PMFG) relaxes the sparsity of the MST while preserving planarity [17]. For N nodes, a PMFG contains $3N - 6$ edges and includes the MST as a subgraph when both are constructed from the same ranked distance matrix. Because the PMFG can retain triangles and 4-cliques, it is more suitable for studying local connectivity, clustering, hub reorganization and group-level structure. Tumminello et al. [18] further explored the relationship between correlation, hierarchies and networks in financial markets. Massara et al. [26] later introduced the Triangulated Maximally Filtered Graph (TMFG), which extends this class of planar filtering methods and improves scalability for large systems.

Filtered networks are relevant for risk analysis. Pozzi et al. [19] used PMFG-based analysis to show that peripheral assets in financial networks can offer better diversification properties than central hubs. Aste et al. [20] studied correlation dynamics in volatile markets using filtered network approaches. Kenett et al. [27] used partial correlation analysis to reveal the dominant role of the financial sector in stock market networks. These studies motivate the hub-reorganization analysis presented here. By comparing original and RMT-filtered networks, we ask whether centrality is driven mainly by market-wide co-movement or whether localized central nodes emerge after removing the dominant market mode.

2.4. Financial networks in emerging markets

Network methods have also been applied to emerging markets, where market structure can differ substantially from that of large developed equity markets. Emerging markets may combine concentrated sectoral composition, macroeconomic sensitivity, lower liquidity in parts of the market and exposure to global risk cycles. These features make them attractive systems for dependency analysis because the correlation structure may reflect both domestic and international channels.

Tabak et al. [7] studied the structure of the Brazilian interbank network, documenting structural properties and vulnerability to targeted failures. Silva et al. [8] analysed the structure and dynamics of the global financial network, with attention to how emerging economies connect to the broader system. Wang et al. [28] studied the correlation structure of world stock markets using Pearson and partial-correlation networks, showing that emerging markets can display distinctive dependency patterns.

Studies focused more directly on Brazilian markets include Pereira et al. [10], who applied detrended cross-correlation analysis to construct multiscale networks of stock markets including B3, and Leonel Caetano and Yoneyama [11], who applied complex network analysis to the Brazilian stock market. Stosic et al. [9] studied correlations and cross-correlations in Brazilian agrarian commodities and stocks. These studies show that Brazilian financial data contain rich structural properties, but the literature remains fragmented across interbank networks, commodities, cross-market relations and equity-market structure.

The present paper contributes to this literature by focusing on a broad B3 equity universe and by connecting spectral filtering, planar networks and volatility forecasting in a single design. This integration matters because a network's interpretation depends on the matrix used to build it. A network constructed from the original correlation matrix may be dominated by the market mode, while a network constructed from group-mode components may reveal weaker but more localized dependency channels. This distinction is especially relevant in emerging markets, where common macroeconomic exposure and sectoral heterogeneity coexist.

2.5. Volatility forecasting and machine learning

Volatility forecasting has increasingly incorporated machine-learning methods because realized-volatility targets can be combined with large feature sets, nonlinear interactions and strict temporal validation. The theoretical foundations of realized-volatility measurement were established by Andersen et al. [29], and Poon and Granger [30] give an influential review of volatility forecasting methods. In the present study, the forecasting component is empirical and predictive: it asks whether features extracted from RMT and filtered networks add information to supervised models for future realized volatility.

Forecast evaluation is delicate because realized volatility is an observable proxy for latent conditional variance. Loss functions can rank models differently depending on their robustness to measurement error and their treatment of volatility underestimation. QLIKE is widely used because of its robustness properties under noisy volatility proxies [31]. For this reason, the forecasting experiment in this paper evaluates models using multiple metrics, including MAE, RMSE, QLIKE and out-of-sample R^2 .

Machine-learning methods offer flexible nonlinear approximations and can incorporate large feature sets. Random Forests [32], gradient boosting [33] and regularized linear models [34] can combine lagged volatility, rolling statistics, market-wide variables and network centrality measures. Bucci [35] applied neural networks to realized-volatility forecasting, showing that flexible models can capture nonlinear volatility dynamics. Christensen et al. [36] developed a machine-learning framework for volatility forecasting and emphasized the importance of careful feature engineering. Zhang and Broadstock [37] explored how financial network structure relates to market risk from a machine-learning perspective.

In financial forecasting, however, flexibility can amplify noise and produce unstable gains if temporal validation is not strict. For this reason, we use a chronological train-validation-test split and treat network features as potential additional predictors, not replacements for lagged volatility information. The nested feature-set design follows an ablation logic: Set A measures the predictive value of standard volatility information, Set B tests the added value of market and RMT variables, and Set C tests whether MST and PMFG structure adds further information beyond these baseline predictors. Simple MLP and CNN-1D models are included as incremental neural extensions rather than as an exhaustive architecture search.

3. Data and asset universe

3.1. BovDB/B3 data source

The dataset consists of Brazilian equity prices from BovDB and its extended version BovDBv2 [38, 39]. BovDB was originally proposed as an open dataset of stock prices for companies traded on the Brazilian Stock Exchange (B3), covering the period from 1995 to 2020 and applying cumulative adjustment factors to correct prices for corporate events. BovDBv2 is distributed through Harvard Dataverse as the *Dataset for Intraday Analysis of B3 stock prices*, containing daily price information from 1995 to 2024 and intraday price data for Brazilian stocks from January 1995 to July 2024.

Although BovDBv2 includes intraday information, we use daily adjusted prices because the objective is to study long-run cross-sectional dependencies, Random Matrix Theory filtering and network structure over a broad historical window. Daily data give a stable sampling frequency for estimating correlations across many assets and reduce microstructure effects that would be more prominent in intraday analysis. The main objects of interest are adjusted closing prices, logarithmic returns and correlation matrices constructed from synchronized daily observations.

3.2. Adjusted prices

Adjusted closing prices are used because long historical equity samples are affected by stock splits, reverse splits, dividends, interest on equity, bonuses, corporate reorganizations and share-class changes. Without adjustment, these events can generate artificial jumps that contaminate return distributions, correlations and network structure. Return construction, standardization and correlation estimation are defined formally in Section 4.

3.3. Asset filters and equity universe definition

The core equity universe is obtained through deterministic filters applied to the BovDB/B3 daily database. The starting point is the set of B3 instruments with daily adjusted price information. We then exclude instruments that are not direct listed equities, including Brazilian Depositary Receipts (BDRs), exchange-traded funds (ETFs), investment funds, benchmark indices and other non-equity vehicles. This filter keeps the dependency matrix focused on listed equity shares instead of mixing equities with instruments that mechanically track baskets, benchmarks or foreign securities.

After this class filter, we retain observations only when adjusted closing prices are valid and strictly positive. Daily logarithmic returns are computed from adjusted closing prices, and extreme adjusted-return observations are inspected as part of the data-quality procedure. These observations are used as exclusion criteria only when they indicate clear corporate-action adjustment inconsistencies. Assets are then required to satisfy a minimum historical-coverage condition, and the remaining assets are synchronized on common trading dates. The formal notation for this synchronized return panel is given in Section 4.1.

3.4. Core historical universe

The core historical universe contains 58 assets with sufficient data coverage for the main RMT, clustering and network analysis. Table 1 summarizes this core universe and the smaller representative-asset sample. The universe is historical, not purely current. It allows us to study a long Brazilian equity sample while keeping the cross-section large enough for network analysis and small enough for interpretable visualizations.

Table 1: Summary of the samples used in this study.

Sample	Assets	Period	Main use
Representative assets	3	2006–2025	Stylized facts and initial forecasting experiment
Core historical universe	58	1998–2025	Correlation, RMT, clustering and networks

The main spectral analysis uses a complete synchronized panel with $N = 58$ assets and $T = 1527$ common daily observations. The complete-case restriction is important because the empirical correlation matrix must be estimated from a common set of trading dates across all assets. This avoids pairwise changes in sample composition and ensures that eigenvalues, eigenvectors, distances and network edges are derived from a single common correlation matrix.

3.5. Demonstration and forecasting assets

We use PETR4, VALE3 and BBDC4 as representative assets for stylized facts and as the initial forecasting targets. They were selected because they are large, liquid and economically distinct Brazilian equities. PETR4 represents oil and gas exposure, with sensitivity to commodity prices and firm-specific institutional events. VALE3 represents basic materials and global mining exposure, with sensitivity to commodity cycles and exchange-rate dynamics. BBDC4 represents the financial sector and domestic credit conditions.

This small panel is not intended to represent the entire Brazilian equity market. It offers a compact and interpretable entry point into the analysis, illustrating heavy tails, volatility clustering, return extremes, absolute-return persistence and the feasibility of the forecasting design.

3.6. Sector and subsector classification

Each asset in the core historical universe is assigned to a macro-sector, subsector and segment classification. The macro sectors include Financials, Oil Gas and Biofuels, Basic Materials, Utilities, Industrials, Consumer Cyclical, Consumer Non-Cyclical, Real Estate and Telecommunications. These labels are used for descriptive comparison, sectoral correlation tests, heatmap interpretation and subsector-level aggregation.

Utilities and Real Estate are kept separate because their risk drivers differ substantially. Utilities are more directly linked to regulation, concessions, inflation adjustment mechanisms and defensive cash flows. Real Estate is more sensitive to credit conditions, interest rates and the real-estate cycle. Within Basic Materials, the subsector classification distinguishes Mining, Steel and Metallurgy, Chemicals, and Paper and Cellulose. This distinction matters in Brazil because commodity-linked firms can share global exposure while still belonging to different industrial chains.

3.7. Data quality notes and limitations

The main data-quality issue is the distinction between genuine extreme returns and artificial jumps caused by corporate-action adjustment problems. Adjusted prices reduce the impact of splits, dividends and related events, but they do not guarantee that every historical corporate action is perfectly corrected. This limitation is especially relevant in long emerging-market samples, where ticker histories, share-class changes and corporate reorganizations can be complex.

For PETR4, VALE3 and BBDC4, the largest daily returns observed in the 2006–2025 window match stress periods such as the 2008 global financial crisis, the Brazilian recession period around 2015–2016 and the 2020 COVID shock. Extreme-return validation is treated as a separate quality-control step, and assets with implausible corporate-action-related jumps are excluded from the main analyses.

The main results are based on the core historical universe and on synchronized daily returns. This choice yields a common synchronized panel for estimating correlation matrices, eigenvalue spectra and network structure, but it also imposes limits. Results may differ under intraday sampling, alternative liquidity filters, rolling-window universes, or broader target sets for forecasting. These extensions are left for robustness analysis and future work.

4. Methodology

The methodology transforms adjusted B3 prices into return matrices, correlation matrices, spectral components, filtered financial networks and volatility-forecasting features. Figure 1 summarizes the workflow. The workflow estimates dependencies and separates market-wide motion, group-level structure and noise-compatible components before evaluating whether the resulting network structure contains predictive information.




Figure 1: Overview of the empirical workflow. Adjusted daily B3 prices are filtered to define the equity universe and construct synchronized return panels. The workflow proceeds through stylized-fact checks, correlation estimation, Random Matrix Theory filtering, hierarchical clustering, MST/PMFG network construction, subsector aggregation and volatility forecasting with machine-learning, deep-learning, RMT and network-based features.

4.1. Notation and return matrix

We fix notation before defining the main steps. The index $i = 1, \dots, N$ identifies assets, and the index $t = 1, \dots, T$ identifies trading days in a synchronized panel. In the main spectral, clustering and network analysis, $N = 58$ is the number of assets and $T = 1527$ is the number of common daily observations after synchronization. Thus T is the number of dates in the complete panel, not the total number of scalar return entries.

Let $P_{i,t}^{\text{adj}}$ denote the adjusted closing price of asset i on trading day t . Daily logarithmic returns are computed as

$$r_{i,t} = \log P_{i,t}^{\text{adj}} - \log P_{i,t-1}^{\text{adj}}, \quad (1)$$

where $P_{i,t}^{\text{adj}} > 0$ and $P_{i,t-1}^{\text{adj}} > 0$. Log returns are used because they are additive across time and standard in studies of financial returns.

The return panel is represented by the matrix

$$\mathbf{R} = (r_{i,t})_{t=1,\dots,T; i=1,\dots,N} \in \mathbb{R}^{T \times N}. \quad (2)$$

For correlation-based analysis, each return series is standardized within the relevant sample:

$$\tilde{r}_{i,t} = \frac{r_{i,t} - \mu_i}{\sigma_i}, \quad (3)$$

where μ_i and σ_i are the sample mean and standard deviation of asset i . The standardized return matrix is denoted by $\tilde{\mathbf{R}}$.

Let $\mathbf{C}$ denote the Pearson correlation matrix, and let C_{ij} be the Pearson correlation between assets i and j . In the Random Matrix Theory step, λ_k denotes the k -th eigenvalue of $\mathbf{C}$, $\mathbf{u}_k$ denotes its normalized eigenvector and $Q = T/N$ denotes the sample-size-to-dimension ratio used in the Marčenko–Pastur benchmark.

4.2. Asset filtering and complete return panel

Let $\mathcal{A}_0$ denote the initial set of B3 instruments available in the BovDB/B3 daily database. The data section explains the economic motivation for excluding non-direct-equity instruments. Formally, let $\mathcal{X}$ be the excluded set containing BDRs, ETFs, investment funds, benchmark indices and other non-equity vehicles. The eligible direct-equity set is

$$\mathcal{A}_1 = \mathcal{A}_0 \setminus \mathcal{X}, \quad (4)$$

where each element of $\mathcal{A}_1$ is an asset identifier, not an observation date.

For each remaining asset, observations are retained only when adjusted closing prices are valid:

$$P_{i,t}^{\text{adj}} > 0. \quad (5)$$

Daily log returns are then computed from adjusted prices. Assets with insufficient historical coverage are removed according to

$$T_i \geq T_{\min}, \quad (6)$$

where T_i is the number of valid return observations for asset i and $T_{\min}$ is the minimum coverage threshold required for the target analysis.

After filtering, assets are synchronized by common trading dates. The RMT, clustering and network analyses are performed on a complete return panel

$$\mathbf{R}_{\text{core}} \in \mathbb{R}^{T \times N}, \quad (7)$$

with $N = 58$ assets and $T = 1527$ complete daily observations in the main spectral analysis. Complete-panel estimation ensures that all eigenvalues, eigenvectors, distances and network edges are derived from a single common correlation matrix estimated from assets observed on common trading dates.

4.3. Stylized facts and descriptive diagnostics

We first document stylized facts for representative B3 assets. For each asset, we compute the sample mean, standard deviation, skewness, excess kurtosis, minimum and maximum returns, value-at-risk at 1% and 5%, counts of large absolute returns and autocorrelations of absolute returns.

For an asset return series $\{r_t\}$, the autocorrelation of absolute returns at lag ℓ is

$$\rho_{|r|}(\ell) = \text{Corr}(|r_t|, |r_{t-\ell}|). \quad (8)$$

These checks assess whether B3 returns display the regularities commonly associated with financial time series: heavy tails, volatility clustering and persistent volatility [2].

4.4. Correlation matrix estimation

The Pearson correlation matrix is estimated from standardized returns as

$$C_{ij} = \frac{1}{T-1} \sum_{t=1}^T \tilde{r}_{i,t} \tilde{r}_{j,t}. \quad (9)$$

Equivalently, in matrix form,

$$\mathbf{C} = \frac{1}{T-1} \tilde{\mathbf{R}}^\top \tilde{\mathbf{R}}. \quad (10)$$

The matrix $\mathbf{C} \in \mathbb{R}^{N \times N}$ is symmetric and has unit diagonal. For pairwise descriptive analysis, correlations may be computed with available observations for each pair. For RMT, clustering, MST and PMFG construction, the complete synchronized panel is used to obtain a single consistent $N \times N$ matrix.

4.5. Sectoral correlation comparison

Each asset is assigned to a macro-sector and subsector. The sectoral comparison asks a simple question: are correlations between assets in the same macro-sector larger than correlations between assets in different macro-sectors? To answer this, we classify every unordered asset pair (i, j) into one of two groups. The indicator s_{ij} is not a correlation coefficient; it only records whether the two assets belong to the same macro-sector:

$$s_{ij} = \begin{cases} 1, & \text{if assets } i \text{ and } j \text{ belong to the same macro-sector,} \\ 0, & \text{otherwise.} \end{cases} \quad (11)$$

After this classification, the corresponding Pearson correlations C_{ij} are split into two samples. Pairs with $s_{ij} = 1$ form the within-sector correlation distribution, while pairs with $s_{ij} = 0$ form the between-sector correlation distribution. We then compare the two distributions using mean, median, dispersion and a one-sided Mann–Whitney U test. This step evaluates whether the economic sector classification is reflected in the dependency matrix before any RMT filtering is applied.

4.6. Rolling market synchronization

Static full-sample correlations hide time variation in market dependence. To measure market synchronization, rolling pairwise correlations are estimated over a window of length w . For each window ending at time t , the average market correlation is

$$\bar{\rho}_t^{(w)} = \frac{2}{N(N-1)} \sum_{1 \leq i < j \leq N} \rho_{ij,t}^{(w)}, \quad (12)$$

where $\rho_{ij,t}^{(w)}$ is the Pearson correlation between assets i and j computed within the rolling window. Elevated values of $\bar{\rho}_t^{(w)}$ indicate stronger market-wide synchronization and motivate the subsequent spectral decomposition.

4.7. Random Matrix Theory filtering

Random Matrix Theory is used to distinguish structured collective modes from eigenvalue fluctuations compatible with sampling noise. Let the spectral decomposition of the empirical correlation matrix be

$$\mathbf{C} = \sum_{k=1}^N \lambda_k \mathbf{u}_k \mathbf{u}_k^\top, \quad (13)$$

where λ_k is the k -th eigenvalue and $\mathbf{u}_k$ is the corresponding normalized eigenvector. Eigenvalues are sorted in decreasing order, so $k = 1$ denotes the largest mode.

For a random correlation matrix estimated from N independent standardized time series of length T , the Marčenko–Pastur bounds are

$$\lambda_{\pm} = \sigma^2 \left(1 + \frac{1}{Q} \pm 2\sqrt{\frac{1}{Q}} \right), \quad Q = \frac{T}{N}. \quad (14)$$

Since standardized returns are used, $\sigma^2 = 1$. Eigenvalues inside $[\lambda_-, \lambda_+]$ are treated as compatible with random-matrix noise, while eigenvalues above λ_+ are candidates for collective modes.

Let K denote the number of sample eigenvalues above λ_+ . The dominant market component is defined as

$$\mathbf{C}^{\text{market}} = \lambda_1 \mathbf{u}_1 \mathbf{u}_1^\top. \quad (15)$$

The group or sector component is reconstructed from the remaining retained eigenvalues:

$$\mathbf{C}^{\text{group}} = \sum_{k=2}^K \lambda_k \mathbf{u}_k \mathbf{u}_k^\top. \quad (16)$$

The retained-mode matrix is

$$\mathbf{C}^{\text{filtered}} = \sum_{k=1}^K \lambda_k \mathbf{u}_k \mathbf{u}_k^\top, \quad (17)$$

and the residual noisy component is

$$\mathbf{C}^{\text{noise}} = \sum_{k=K+1}^N \lambda_k \mathbf{u}_k \mathbf{u}_k^\top. \quad (18)$$

Note that partial spectral reconstructions generally do not preserve the unit diagonal of the original correlation matrix. When these filtered matrices are used for downstream tasks, they are rescaled to unit diagonal if the subsequent method requires valid Pearson correlation bounds. This decomposition separates networks dominated by broad market co-movement from networks that emphasize localized group and sectoral structure.

4.8. Mantegna distance and hierarchical clustering

Correlations are converted into distances using the metric introduced by Mantegna [16]:

$$d_{ij} = \sqrt{2(1 - C_{ij})}. \quad (19)$$

For a valid correlation matrix, this transformation maps strongly positively correlated assets to short distances and less positively correlated or negatively correlated assets to longer distances. The resulting distance matrix is used for hierarchical clustering, ordered heatmaps, MST and PMFG construction.

Hierarchical clustering is applied to the distance matrix using average linkage (UPGMA). The cophenetic correlation [40] is used to assess how faithfully the dendrogram preserves pairwise distances. Ordered heatmaps use the dendrogram ordering to rearrange the correlation matrices, making block structures visually identifiable. Comparing original, filtered and group-mode heatmaps helps distinguish broad market co-movement from localized sectoral dependency.

4.9. Minimum Spanning Tree and Planar Maximally Filtered Graph

A financial network is represented as

$$G = (V, E, w), \quad (20)$$

where V is the set of assets, E is the retained set of edges and $w_{ij} = C_{ij}$ is the dependency weight associated with edge (i, j) .

The Minimum Spanning Tree is the connected acyclic graph that minimizes total distance:

$$\text{MST} = \arg \min_{E'} \sum_{(i,j) \in E'} d_{ij}, \quad (21)$$

subject to (V, E') being connected and acyclic. For N nodes, the MST has $N - 1$ edges. With $N = 58$, this gives 57 edges.

The Planar Maximally Filtered Graph is constructed by considering candidate edges in increasing order of distance, retaining them if they do not violate planarity. For N nodes, a PMFG has

$$|E_{\text{PMFG}}| = 3N - 6 \quad (22)$$

edges, which gives 168 edges for $N = 58$. Unlike the MST, the PMFG can contain cycles, triangles and 4-cliques. Under the same ranking of pairwise distances, the MST is contained in the PMFG as a subgraph. This makes the PMFG a less aggressive graph filter that preserves richer local geometry.

For each network, we report edge correlations, same-sector and same-subsector edge ratios, average clustering, node degree, betweenness centrality and hub rankings. Comparing original and group-mode networks identifies whether centrality is driven by broad market co-movement or by localized dependency channels.

4.10. Aggregated subsector dependency network

The asset-level networks are complemented by an aggregated subsector dependency network. Let $g(i)$ denote the subsector associated with asset i . For two subsectors a and b , the aggregated dependency weight is defined as

$$W_{ab} = \frac{1}{|\mathcal{E}_{ab}|} \sum_{(i,j) \in \mathcal{E}_{ab}} w_{ij}, \quad (23)$$

where

$$\mathcal{E}_{ab} = \{(i, j) \in E : g(i) = a, g(j) = b\}. \quad (24)$$

The aggregation converts ticker-level dependencies into a subsector-level map of dependency channels. The resulting network is summarized using edge density, average dependency, weighted degree and betweenness centrality.

4.11. Volatility forecasting formulation

The forecasting task is formulated as a supervised time-series prediction problem. For target asset i and horizon h , realized volatility is defined as

$$RV_{i,t,h} = \sqrt{\sum_{\tau=1}^h r_{i,t+\tau}^2}. \quad (25)$$

The objective is to predict $RV_{i,t,h}$ using an information set $\mathcal{F}_t$ available at time t :

$$\widehat{RV}_{i,t,h} = f_{\theta}(\mathcal{F}_t). \quad (26)$$

For forecast evaluation under metrics such as QLIKE, these realized volatility and predicted volatility measures are squared to represent realized variance proxies. The target horizons are $h = 5$ and $h = 20$ trading days. The initial forecasting experiment uses PETR4, VALE3 and BBDC4, with a strict temporal split into training, validation and test periods to reduce look-ahead bias and preserve the chronological ordering of the forecasting experiment.

4.12. Machine-learning and deep-learning feature sets

Three feature sets are used in a nested ablation design. Set A contains classical return and volatility predictors, including rolling volatilities, rolling absolute returns, lagged realized volatility, EWMA volatility and higher-moment rolling statistics. Set B adds market and RMT variables such as the largest eigenvalue, market-mode trace share and number of eigenvalues above the random-matrix bound. Set C adds MST and PMFG features, including degree, betweenness and clustering measures from original and group-mode networks.

This nesting defines a controlled ablation:

$$\mathcal{F}_t^A \subset \mathcal{F}_t^B \subset \mathcal{F}_t^C. \quad (27)$$

Set A measures the predictive value of standard lagged return and volatility information alone. Set B tests the added value of market-wide and RMT features. Set C tests whether network structure adds further information beyond these

baselines. Ridge regression, Random Forest and histogram gradient boosting are used to evaluate these increasingly rich information sets [34, 32, 33].

The deep-learning extension is intentionally incremental. A multilayer perceptron (MLP) is trained on the same tabular feature sets, using positive-output regression for realized volatility. A one-dimensional convolutional neural network (CNN-1D) is trained on short rolling windows of return and volatility features. These neural models are included to test whether a simple nonlinear representation adds forecasting value; they are not presented as an exhaustive search over neural architectures [35, 36].

4.13. Forecast evaluation metrics

Forecasts are evaluated using MAE, RMSE, out-of-sample R^2 and QLIKE. For realized values y_t and forecasts $\hat{y}_t$, the MAE and RMSE are

$$\text{MAE} = \frac{1}{M} \sum_{t=1}^M |y_t - \hat{y}_t|, \quad (28)$$

and

$$\text{RMSE} = \sqrt{\frac{1}{M} \sum_{t=1}^M (y_t - \hat{y}_t)^2}. \quad (29)$$

The out-of-sample coefficient of determination is

$$R_{\text{oos}}^2 = 1 - \frac{\sum_{t=1}^M (y_t - \hat{y}_t)^2}{\sum_{t=1}^M (y_t - \bar{y}_{\text{train}})^2}. \quad (30)$$

The QLIKE loss is

$$\text{QLIKE}_t = \log(\hat{\sigma}_t^2) + \frac{\sigma_t^2}{\hat{\sigma}_t^2}, \quad (31)$$

where $\hat{\sigma}_t^2$ is the predicted variance proxy and σ_t^2 is the realized variance proxy. QLIKE is included because it is widely used for volatility forecast comparison when realized volatility is an imperfect proxy for latent variance [31].

5. Stylized facts and correlation structure

5.1. Stylized facts of selected B3 assets

Figure 2 summarises the return properties of PETR4, VALE3 and BBDC4. We use these assets as representative examples and as the initial forecasting targets because they are liquid, economically important and sectorally distinct; they are not intended to be a proxy for the entire Brazilian equity market. The figure combines four checks that are standard in the econophysics literature [2]: normalized prices, daily returns, the complementary cumulative distribution of absolute returns and the autocorrelation of absolute returns.




Figure 2: Stylized facts for PETR4, VALE3 and BBDC4 over 2006–2025. The panels show normalized prices, daily log returns, the complementary cumulative distribution of absolute returns and autocorrelations of absolute returns. The return panel uses three vertically aligned subpanels with a common scale and a zero reference line for each asset.

The normalized-price panel illustrates heterogeneous long-run trajectories across the three assets. This heterogeneity is expected because the selected firms belong to distinct economic segments and are exposed to different sectoral and macro-financial drivers. The return panel displays periods in which large movements cluster over time, a pattern indicating volatility clustering. The CCDF panel suggests that large absolute returns occur more frequently than would be expected under a Gaussian benchmark, pointing to heavy-tailed returns. The autocorrelation panel shows positive dependence in $|r_t|$ over several lags, indicating persistent volatility [2].

5.2. Descriptive statistics

Table 2 quantifies the same diagnostics. All three series have means close to zero, substantial kurtosis and non-negligible counts of large absolute returns. PETR4 has the largest negative daily observation in the 2006–2025 window, while VALE3 and BBDC4 also show stress-period extremes. The autocorrelations of absolute returns remain positive across several lags, especially for VALE3.

The descriptive statistics match the visual reading of Figure 2. The return distributions display substantial departures from a Gaussian benchmark, including excess kurtosis and persistent absolute-return autocorrelations. These properties match the well-documented stylized facts of financial returns [2] and support treating the Brazilian market as a complex system in which dependency structure, not marginal normality, is the relevant modelling object.

Table 2: Descriptive statistics for the three demonstration assets, 2006–2025.

Symbol	n	Mean	Std.	Skew.	Ex. kurt.	Min.	Max.	Var $R_{1\%}$	Var $R_{5\%}$	$ r > 10\%$	$ r > 20\%$	ACF $_1$	ACF $_5$	ACF $_{20}$	ACF $_{60}$
BBDC4	4953	0.00028	0.02180	0.02	7.96	-0.1910	0.1999	-0.0559	-0.0312	16	0	0.207	0.174	0.122	0.044
PETR4	4953	0.00041	0.02755	-0.76	11.55	-0.3524	0.2007	-0.0767	-0.0408	41	4	0.254	0.209	0.140	0.062
VALE3	4953	0.00045	0.02576	-0.22	7.51	-0.2814	0.1936	-0.0676	-0.0381	26	2	0.231	0.207	0.169	0.119

Notes: Mean and Std. are the sample mean and standard deviation of daily log returns. Skew. is sample skewness and Ex. kurt. is excess kurtosis relative to the Gaussian benchmark. $\text{VaR}_{1\%}$ and $\text{VaR}_{5\%}$ are the empirical 1% and 5% lower-tail return quantiles. $|r| > 10\%$ and $|r| > 20\%$ count daily observations whose absolute log return exceeds the stated threshold. ACF_ℓ is the autocorrelation of absolute returns at lag ℓ trading days.

5.3. Static correlation distribution

We now turn from individual return behaviour to cross-sectional dependence. The primary object is the complete correlation matrix for the 58-asset synchronized panel. The distributional summaries below compress this matrix into pairwise statistics before the RMT decomposition is discussed. The core historical universe contains 58 assets, producing $N(N-1)/2 = 1653$ unique asset pairs. Figure 3 summarizes the full pairwise correlation distribution and then splits pairs into within-sector and between-sector groups.




Figure 3: Correlation structure in the core historical universe. Panel (a) shows the full Pearson correlation distribution. Panel (b) compares within-sector and between-sector correlations using the initial macro-sector classification.

The full distribution is centred in a moderate positive range, with mean correlation close to 0.24 and median correlation close to 0.23. This level of average dependence suggests the presence of a common component, but the distribution remains dispersed. The dispersion is important because it leaves room for sectoral and subsectoral organization instead of a single homogeneous market block.

5.4. Sectoral correlation comparison

The sector split gives a descriptive check of this interpretation. Within-sector correlations are visibly shifted to the right relative to between-sector correlations. Table 3 reports the corresponding pair counts and summary statistics. These results indicate that the initial macro-sector map is reflected in the dependency matrix.

Table 3: Within-sector and between-sector correlation summary.

Group	Pairs	Mean	Median	Std. dev.	Min	Max
Within-sector	275	0.3433	0.3082	0.1819	0.0534	0.9402
Between-sector	1378	0.2249	0.2076	0.1166	-0.0236	0.5552

The difference is economically meaningful. Within-sector pairs have mean correlation 0.3433, while between-sector pairs have mean correlation 0.2249. The one-sided Mann–Whitney test for the alternative that within-sector correlations are larger gives $p = 1.33389 \times 10^{-24}$. Although the Mann–Whitney test assumes independent observations—an assumption violated by overlapping pairwise correlations—the large effect size supports the descriptive interpretation that sectoral groupings are reflected in the dependency matrix.

5.5. Dynamic market synchronization

Full-sample correlations are limited because market dependence changes over time. Figure 4 shows the rolling average market correlation. This statistic compresses the correlation matrix into a time series that measures how synchronized the market is at each point in time.




Figure 4: Rolling average market correlation for the core historical universe.

The rolling average correlation tends to rise around major stress periods, including the 2008 global financial crisis, the Brazilian recession period around 2015–2016 and the 2020 COVID shock. These episodes indicate periods of stronger market-wide synchronization, which helps explain the dominant leading eigenvalue in the full-sample correlation matrix.

Figure 5 adds a more granular view by showing selected pairwise dynamic correlations under different estimators. We use these estimates to show that individual dependencies can change substantially over time, not to replace the full matrix.

The dynamic pairwise estimates illustrate why a single full-sample correlation requires care. Short-window correlations react quickly but are noisy; longer-window estimates are smoother but slower; EWMA estimates give an intermediate view. This descriptive evidence motivates the subsequent use of spectral filtering and networks as structured summaries of a noisy and time-varying dependency system.




Figure 5: Dynamic pairwise correlations for selected asset pairs. Short-window, longer-window and EWMA estimates highlight the instability of individual dependencies.

6. Random Matrix filtering

6.1. Empirical eigenspectrum

Random Matrix Theory supplies the spectral benchmark for distinguishing structured collective modes from sampling noise in the empirical correlation matrix [13, 14]. The core historical matrix uses $N = 58$ assets and $T = 1527$ observations, giving $Q = T/N \approx 26.33$. The Marčenko–Pastur bounds are $\lambda_- = 0.6482$ and $\lambda_+ = 1.4278$.

Table 4: RMT summary for the core historical universe.

N	T	Q	λ_-	λ_+	λ_1	λ_2	λ_3	λ_4	λ_5	Above λ_+
58	1527	26.33	0.6482	1.4278	21.6505	3.0957	1.7839	1.6543	1.4835	5

Table 4 reports the numerical inputs for RMT filtering. With $N = 58$ assets and $T = 1527$ complete observations, the aspect ratio is $Q = 26.33$, which implies a relatively narrow Marčenko–Pastur noise band from $\lambda_- = 0.6482$ to $\lambda_+ = 1.4278$. The first eigenvalue, $\lambda_1 = 21.6505$, is therefore not a marginal deviation from the random-matrix benchmark but a dominant collective component. The next four eigenvalues, $\lambda_2 = 3.0957$, $\lambda_3 = 1.7839$, $\lambda_4 = 1.6543$ and $\lambda_5 = 1.4835$, also exceed the upper random bound. There are 5 sample eigenvalues outside the band: these are the candidate collective modes, indicating one broad market mode and four additional modes that may be associated with group, sectoral or localized dependency structure.

Figure 6 shows the spectrum against the random benchmark. The largest eigenvalue lies far above the Marčenko–Pastur upper bound. Four additional eigenvalues also exceed the upper edge, indicating a dominant market mode and a small number of group or sector modes.




Figure 6: Empirical eigenvalue spectrum and Marčenko–Pastur bounds. The largest eigenvalue is far outside the random band and captures the market-mode component.

6.2. Market Mode and significant eigenvalues

The largest eigenvalue is 21.6505, corresponding to approximately 37.3% of the total trace of the correlation matrix. Its eigenvector has broadly positive loadings, confirming its interpretation as a market mode. It captures broad co-movement and explains why the original correlation matrix has a strong positive common component. The next four eigenvalues are 3.0957, 1.7839, 1.6543 and 1.4835. These values are above λ_+ and act as candidates for group or sector structure.

The market mode is meaningful but visually dominant. If a network is built directly from the original matrix, many edges can reflect market-wide co-movement instead of local structure. We construct the group-mode matrix to ask what remains after the broad common component is removed.

6.3. Matrix reconstruction

The reconstruction step checks that the spectral decomposition and the filtered matrices are being computed accurately. If all eigenvalue–eigenvector modes are recombined, the original correlation matrix is recovered up to

floating-point precision. Table 5 reports this numerical check: the Frobenius reconstruction error is approximately 2.74×10^{-14} , and the maximum absolute reconstruction error is approximately 3.99×10^{-15} . The reported unit diagonal refers to the filtered matrix after the partial reconstruction has been rescaled for downstream correlation-based distances. This check reduces the possibility that subsequent network differences are driven by numerical reconstruction error rather than by the intended selection of market, group and noise-compatible modes.

Table 5: RMT reconstruction diagnostics.

Diagnostic	Value
Frobenius reconstruction error	2.74×10^{-14}
Maximum absolute reconstruction error	3.99×10^{-15}
C_{filtered} diagonal min	1.0000
C_{filtered} diagonal max	1.0000
Eigenvalues above λ_+	5
Market-mode trace share	0.3733

Figure 7 shows the resulting matrix components. The original-correlation panel is the complete empirical correlation matrix used throughout the paper. The market-mode matrix captures broad positive co-movement. The group-mode matrix displays weaker localized patterns after the market mode is removed. The noise component contains the random-band contribution, and the filtered matrix recombines the retained modes.

Empirically, the eigenvectors beyond the leading component show loading patterns suggesting localized or sector-related dependencies. This residual organization motivates the clustering and network sections: if localized dependency patterns remain after market-mode removal, MST and PMFG structures should differ when constructed from the group-mode matrix.



Figure 7: Empirical and RMT-filtered matrix components. The original-correlation panel shows the complete 58-asset correlation matrix; the remaining panels separate it into market, group/sector, noise and filtered components.

RMT-filtered networks for B3 volatility forecasting




Figure 8: Top eigenvector loadings for the leading RMT modes. Eigenvector 1 captures the market-mode component; subsequent eigenvectors are read as candidate group or localized sectoral components.

6.4. Top eigenvector loadings

Eigenvalues identify the strength of collective modes, but eigenvectors help interpret their composition. Figure 8 reports the largest loadings for the leading eigenvectors. The first eigenvector captures the market-mode component because its loadings are broadly positive across the cross-section. We read subsequent eigenvectors as candidate group or localized components.

For each eigenvector, assets are ranked by the magnitude of their loadings. The signs of eigenvector loadings are arbitrary up to multiplication by -1 , so interpretation focuses on relative grouping and magnitude, not sign alone.

The second eigenvector suggests a Basic Materials and commodity-related component. The third contains financial and cross-sector structure. The fourth and fifth eigenvectors are more localized, with Telecommunications and Utilities-like components appearing among the dominant loadings. These labels are economic readings of the loading patterns, not mechanical classifications. Eigenvectors can mix sectors when firms share macroeconomic, liquidity or balance-sheet exposures.

7. Hierarchical clustering and ordered heatmaps

7.1. Ordered heatmaps

RMT decomposes the correlation matrix into spectral components, while ordered heatmaps offer a visual check of whether these components display block-like organization. We order the assets through hierarchical clustering in Figure 9, which displays original, filtered and group-mode structures. The heatmaps serve as diagnostic tools, assessing whether spectral filtering is associated with patterns that can be read in relation to sector and subsector classifications.

The original matrix displays a broad positive background, reflecting the large leading eigenvalue. The filtered matrix preserves part of this organized dependence while reducing the contribution of eigenvalues inside the random-matrix band. The group-mode matrix has lower average correlation by construction, but its ordered representation displays localized block-like patterns that are less dominated by the market-wide component.




Figure 9: Ordered heatmaps for original and RMT-filtered correlation matrices. Hierarchical ordering reveals block structures that are less visible in unordered matrices.

7.2. Dendrograms comparison

Dendrograms give another view of the same distance geometry. Figure 10 compares hierarchical clustering trees for original, filtered and group-mode matrices, and Table 6 reports the corresponding cophenetic correlations. These trees help identify which assets merge at shorter distances and how the hierarchy changes after market-mode removal. We do not formally test branch stability here; the dendrograms are used as exploratory summaries of the distance geometry.

Table 6: Cophenetic correlations for hierarchical clustering.

Matrix	Assets	Cophenetic correlation
Original	58	0.9500
Filtered	58	0.9340
Group Mode	58	0.8705

remains high for the filtered matrix, because removing the leading market is not sufficient to remove the correlation structurally. The lower value indicates

es using Mantegna distances and

aps and dendrograms suggest that
re appears when only the strongest

8. Financial network topology

Clustering summarises hierarchical organisation, but it does not identify graph-theoretic hubs, bridges or local cycles. Financial networks give these additional views by retaining selected high-correlation, short-distance edges under structural constraints. We move from dendrograms to filtered financial networks to examine how dependency structure changes after RMT filtering.

8.1. MST backbone

The MST is the sparsest connected representation of the dependency geometry. With $N = 58$, it contains 57 edges. Figure 11 compares the MST built from the original matrix with the MST built from the group-mode matrix. This comparison illustrates how the MST backbone changes after the leading market component is removed.




Figure 11: Refined MST comparison for original and group-mode correlation matrices. The original MST emphasizes strong market-mode dependencies, while the group-mode MST highlights weaker localized dependency channels.

The original MST has mean edge correlation 0.5796, while the group-mode MST has mean edge correlation 0.1388. This decline is expected because the group-mode network is constructed after excluding the dominant common component. The same-sector edge ratio also changes, from 0.7193 in the original MST to 0.6316 in the group-mode MST. The group-mode MST is not simply a weaker version of the original tree. It represents a different dependency backbone constructed after removing the leading common component.

8.2. PMFG topology

The PMFG is less sparse than the MST because it preserves planarity while retaining $3N - 6 = 168$ edges. It can retain cycles, triangles and 4-cliques, which are absent from the MST. Figure 12 compares original and group-mode PMFGs. This is the main network visualization in the paper because the PMFG retains enough structure to study local connectivity patterns, clustering and hub reorganization.

The original PMFG has mean edge correlation 0.5177. The group-mode PMFG has mean edge correlation 0.1109. Despite this lower average edge strength, the group-mode PMFG has higher average clustering, 0.6653, than the original PMFG, 0.5273. Removing the market mode lowers the absolute level of edge correlations, but the resulting PMFG is associated with stronger local clustering.




Figure 12: Refined PMFG comparison for original and group-mode correlation matrices. The group-mode PMFG reduces market-wide dominance and highlights more localized dependency channels.

8.3. Original vs Group/Sector Mode networks

Table 7 summarizes the main MST and PMFG statistics. Across both graph families, original networks display stronger average edge correlations, while group-mode networks display weaker edge correlations but retain nontrivial sectoral and graph structure. Original networks have stronger average edge correlations, as expected, because they are constructed from matrices that include the leading market component. Group-mode networks have weaker edges but retain patterns aligned with sectoral and subsectoral groupings.

Table 7: Selected MST and PMFG topology statistics.

Network	Matrix	Edges	Mean edge corr.	Same-sector ratio	Clustering	Top degree	Top betweenness	4-cliques
MST	Original	57	0.5796	0.7193	0.0000	RAPT4	ITSA4	0
MST	Group mode	57	0.1388	0.6316	0.0000	ALPA4	ALPA4	0
PMFG	Original	168	0.5177	0.6012	0.5273	BBDC4	GGBR4	55
PMFG	Group mode	168	0.1109	0.5417	0.6653	GUAR3	GUAR3	55

8.4. Network topology metrics

Figure 13 expands the comparison beyond average edge correlation. It summarizes mean edge correlation, same-sector ratio, same-subsector ratio and average clustering across network and matrix types. This figure illustrates that network structure cannot be reduced to a single edge-weight statistic.

The network comparison supports a subsector-level interpretation. Group-mode networks have lower mean edge correlations, while the PMFG clustering coefficient increases. Same-sector and same-subsector ratios remain substantial after market-mode removal. Table 8 confirms that all PMFG variants satisfy the expected planar clique counts. This pattern suggests localized dependency structure, although a formal comparison with graph null models would be required to quantify how unlikely the structure is under randomized alternatives. MST clustering is zero by construction because trees do not contain triangles; the PMFG clustering coefficient is the relevant statistic for local triangular organization.

Table 8: PMFG clique and planarity diagnostics.

Matrix	Nodes	Edges	Triangles	4-cliques	Planar
Original	58	168	166	55	Yes
Filtered	58	168	166	55	Yes
Group Mode	58	168	166	55	Yes

The clique counts are construction checks. For a PMFG with 58 nodes, 168 edges are expected. The presence of triangles and 4-cliques shows that the PMFG retains local graph structures that are absent from the MST. These structures can be transformed into network predictors for the forecasting experiment, although their economic interpretation depends on the underlying edge weights and sector composition.




Figure 13: Network structure comparison across original, filtered and group-mode MST/PMFG representations.

8.5. Hub reorganization after RMT filtering

Hub rankings summarize which assets occupy central positions in the constructed networks. Figure 14 compares hub rankings before and after filtering. In the original PMFG, GGBR4 is the top betweenness node and BBDC4 is the top degree node. In the group-mode PMFG, GUAR3 becomes the leading node by both degree and betweenness.

This change in hub ranking is meaningful for interpretation. Original hubs can reflect market-wide correlation strength and broad exposure to the common component. Group-mode hubs act as localized bridges in the network obtained after removing the leading common component. The appearance of assets such as GUAR3, BBDC3, VALE3, CCRO3, SAPR4 and CMIG3 in group-mode rankings suggests that centrality depends on the matrix used to construct the network. Because hub identities can be sensitive to small changes in edge ranking, we report these rankings as descriptive summaries rather than stable structural labels.

Hub Shift: Impact of RMT Filtering on Network Centrality




Figure 14: Hub-rank comparison across original and group-mode networks. RMT filtering changes which assets occupy central positions in the dependency network.

9. Aggregated subsector dependency network

Asset-level networks can be difficult to interpret when many share classes and sectors are present. The aggregated subsector dependency network converts ticker-level dependencies into a subsector-level map. Nodes represent subsectors and edges summarize average retained dependency strength between subsectors. This is analogous to sector or country aggregation maps used in the financial dependency literature, but adapted to the Brazilian equity universe [3].

We summarize the resulting network in Figure 15 and Table 9. The network contains 22 subsectors and 66 edges, with density 0.2857. The mean edge dependency is 0.0566 and the median edge dependency is 0.0561. These values are lower than typical raw asset-level correlations, which is expected because the network summarizes filtered and aggregated dependencies instead of direct pairwise asset correlations.

The top-degree subsector is Retail, indicating broad direct connectivity in the aggregated dependency map. Electric Utility has the highest betweenness, suggesting that it occupies a bridge-like position in the subsector-level network. Apparel and Footwear has the highest weighted degree, reflecting stronger dependency intensity across its retained connections. These summaries help translate the ticker-level PMFG structure into economic language.

The subsector network also connects the structural results to forecasting. If network structure contains information beyond univariate volatility persistence, then node-level or group-level centrality measures may help explain future volatility regimes. We test this idea in a conservative first experiment.

Table 9: Aggregated subsector network summary.

Subsectors	Edges	Density	Mean edge dep.	Median edge dep.	Top degree	Top betweenness	Top weighted degree
22	66	0.2857	0.0566	0.0561	Retail	Electric Utility	Apparel and Footwear

Notes: Edges are retained subsector-level links in the aggregated dependency network. Density is the ratio between retained and possible undirected subsector links. Mean and median edge dependency are computed from the retained edge weights W_{ab} defined in Section 4.10. Degree counts retained links, weighted degree sums retained edge weights and betweenness is the standard shortest-path betweenness centrality computed on the aggregated graph.




Figure 15: Aggregated subsector dependency network. Nodes represent subsectors, node colour identifies the macro-sector and node size is proportional to the number of assets in the subsector. Edges represent retained subsector-level dependency weights, with thicker edges indicating stronger average dependency.

10. Volatility forecasting

10.1. Forecasting dataset and temporal split

We evaluate whether RMT and network features contain additional information for realized-volatility forecasting. The experiment is deliberately conservative. It uses PETR4, VALE3 and BBDC4 as target assets and compares machine-learning models with simple neural extensions. The data are split temporally: training covers 2006–2017, validation covers 2018–2020 and testing covers 2021–2025. This temporal split is designed to reduce look-ahead bias and to evaluate model selection on observations later than those used for training. Because some network and RMT features are based on static or slowly changing structures, the forecasting experiment is an added-information test, not a fully real-time trading design.

10.2. Machine-learning and deep-learning models

The machine-learning models use three nested feature sets in an ablation design. Set A contains standard lagged return and volatility features. Set B adds market and RMT variables. Set C adds MST and PMFG network variables. This nesting allows the experiment to assess whether network structure adds information beyond lagged-volatility persistence and broad market/RMT features.

The deep-learning extension uses two simple architectures. The MLP is a tabular neural regressor trained on the same feature sets as the machine-learning models. The CNN-1D uses short rolling windows of return and volatility variables. Both neural models use positive-output regression for realized volatility. These models are included as an incremental extension, not as a claim that the neural architecture has been optimized for this market.

10.3. Out-of-sample model comparison

Figure 16 summarizes out-of-sample QLIKE loss, and Table 10 reports the aggregated test-set metrics. Lower QLIKE is better. For each model and horizon, the reported feature set is selected using the validation period and then evaluated on the test period.




Figure 16: Volatility forecasting comparison for machine-learning and deep-learning models. Labels show the validation-selected feature set for each model and horizon.

At the 5-day horizon, Random Forest with Set C obtains the lowest average test QLIKE among the validation-selected models. Ridge, CNN-1D and MLP remain close in MAE and RMSE, but the QLIKE ranking favours the network-augmented Random Forest. This suggests that PMFG and MST variables can add useful information for short-horizon variance-sensitive evaluation, although the differences are descriptive.

At the 20-day horizon, Ridge with Set C obtains the lowest average test QLIKE, while CNN-1D gives competitive MAE, RMSE and out-of-sample R^2 . The MLP has relatively low MAE and RMSE but weaker QLIKE, indicating that it can track average volatility levels while still producing less favourable variance-proxy loss. Overall, the forecasting evidence supports an incremental interpretation: network and RMT features can help, but the gains depend on the model, horizon and evaluation metric.

Table 10: Aggregated test-set forecasting comparison (2021–2025). Values are averaged across PETR4, VALE3 and BBDC4 after validation-based model selection.

Horizon	Model	Feature set	MAE	RMSE	QLIKE	R^2_{oos}
5	Random Forest	Set C (+Network)	0.0161	0.0227	−5.2508	0.1989
5	Ridge	Set C (+Network)	0.0146	0.0214	−5.2303	0.2774
5	HistGradientBoosting	Set C (+Network)	0.0168	0.0234	−5.2228	0.1477
5	CNN-1D	Sequence window	0.0143	0.0213	−5.2098	0.2886
5	MLP	Set A (Classical)	0.0138	0.0210	−5.2095	0.3046
20	Ridge	Set C (+Network)	0.0219	0.0309	−3.8812	0.3803
20	CNN-1D	Sequence window	0.0205	0.0306	−3.8552	0.3951
20	Random Forest	Set C (+Network)	0.0264	0.0368	−3.8455	0.1667
20	HistGradientBoosting	Set C (+Network)	0.0265	0.0371	−3.8417	0.1685
20	MLP	Set B (+Market/RMT)	0.0213	0.0305	−3.4271	0.3726

Notes: MAE is mean absolute error, RMSE is root mean squared error, QLIKE is the variance-proxy loss defined in Section 4.13, and R^2_{OOS} is computed relative to the training-sample mean benchmark. Set A contains standard return and volatility predictors, Set B adds market/RMT variables, and Set C adds MST/PMFG network features.

10.4. Realized vs predicted volatility checks

Aggregate metrics can hide important time-series behaviour. Figure 17 compares realized 20-day volatility with the validation-selected machine-learning or deep-learning forecast for PETR4, VALE3 and BBDC4. We include this plot in the main text because it illustrates how the selected model tracks changing volatility regimes.




Figure 17: Realized versus predicted 20-day volatility for the validation-selected machine-learning or deep-learning model. The vertical dashed line marks the beginning of the 2021–2025 test period.

The plots suggest that broad volatility regimes are tracked reasonably well, but extreme spikes remain difficult to forecast. This matches the aggregate results: flexible models can be competitive, but lagged volatility remains the main source of predictive information. The network variables should therefore be read as complementary structural predictors rather than replacements for time-series volatility information.

10.5. Feature importance and network signal

Random Forest feature importance gives a model-specific check of which predictors contribute most to the fitted model. Figure 18 shows that standard volatility features dominate the model. Rolling absolute returns, lagged realized volatility, rolling volatility and EWMA variables carry most of the importance.




Figure 18: Feature importance for Random Forest volatility forecasting models. Standard volatility predictors dominate the selected rankings, while selected PMFG-based and sector/subsector variables appear with smaller but nonzero importance in the network-augmented feature set.

The strongest features are rolling absolute return over 20 days, lagged realized volatility and longer rolling volatility measures. Network features have smaller importance values than standard volatility predictors, but selected PMFG variables enter the top-ranked predictors, especially at the 20-day horizon. Because Random Forest importance measures can be affected by correlation among predictors, these rankings are model checks, not causal measures of variable relevance. In summary, topology does not replace lagged-volatility features, but it may add structural information when included in a broader feature set.

11. Discussion

Taken together, the results point to a layered interpretation of Brazilian equity dependencies, in line with the multi-scale perspective advocated by Raddant and Di Matteo [3]. The first layer is a market-mode component. It is reflected in the moderate positive average correlation ($\bar{\rho} \approx 0.24$), the increase in rolling average correlation around stress periods and the very large leading eigenvalue ($\lambda_1 = 21.65$, approximately 37.3% of the trace of the correlation matrix). We interpret this component as broad market synchronization. In the Brazilian context, such synchronization may reflect common exposure to domestic monetary conditions, exchange-rate dynamics, country risk and global commodity cycles, although these channels are not separately identified in the present analysis. The component is also methodologically important because it can dominate raw networks and obscure weaker local structures.

The second layer consists of candidate group and sector structure. Within-sector correlations are higher than between-sector correlations on average, and the RMT decomposition identifies four additional eigenvalues above the random upper edge. Eigenvector loadings and group-mode matrices suggest that commodity, financial, telecommunications and utility-related structures contribute to the non-random part of the spectrum. This reading is descriptive, not causal: we identify structure in the dependency matrix, not the economic mechanisms that generate each edge.

The third layer is noise-compatible residual dependence. Not every correlation is economically meaningful, and not every edge retained by a graph is a stable economic relationship. This is why we combine RMT, heatmaps, dendrograms and network metrics. A result is more credible when it appears across several representations, not only in one visualization.

The network results show that market-mode removal changes structure in a systematic way. Original networks have stronger average edge correlations, while group-mode networks have much lower average edge correlations. However, the group-mode PMFG has higher clustering (0.6653 versus 0.5273). Notably, lower average edge strength does not mean absence of structure. It means that the dominant common component has been removed and localized organization becomes easier to observe [17, 18]. The PMFG is well suited for this analysis because it preserves richer graph structure than the MST, retaining 168 edges with 166 triangles and 55 4-cliques that allow hub-rank shifts, clustering changes and subsector dependency maps to be studied in detail.

Clustering summarizes hierarchical organization, but it does not identify graph-theoretic hubs, bridges or local cycles. The transition from dendrograms to filtered financial networks adds information. Hub reorganization after RMT filtering is meaningful: original hubs such as GGBR4 and BBDC4 are central in networks that include the common component, while group-mode hubs such as GUAR3 occupy central positions after the leading component is removed. This suggests that centrality in financial networks depends on the matrix used to construct the network.

The forecasting results require caution. Machine-learning and simple deep-learning models do not produce uniform gains across horizons and metrics. Instead, network-augmented Random Forest models are competitive at the 5-day horizon, while Ridge regression and CNN-1D are competitive at the 20-day horizon. Feature-importance results indicate that lagged volatility variables remain dominant, with selected PMFG variables contributing smaller but nonzero information. The net interpretation is modest: topology adds structural information from the cross-sectional geometry of the market to the temporal dynamics of individual asset variance.

The broader contribution is methodological. We connect econophysics filtering, financial networks and volatility forecasting in a single reproducible workflow for B3 equities. By documenting each step transparently and evaluating network features through a chronological machine-learning and deep-learning design, the workflow offers a template that can be extended to other emerging markets, different time scales and more sophisticated forecasting architectures.

12. Limitations

The first limitation concerns time variation. The RMT and network features used in the initial forecasting exercise are based on full-sample or slowly changing structures. They are useful for understanding the market, but they are not fully real-time predictors unless reconstructed recursively using only information available at each forecast origin. A stricter forecasting design should use rolling RMT and rolling network features.

The second limitation concerns the target set. We use PETR4, VALE3 and BBDC4 in the forecasting experiment. These assets are liquid and economically important, but they do not cover the full B3 cross-section. Future work should extend the forecasting exercise to a larger set of assets and evaluate whether network features help more for some sectors than others.

The third limitation concerns data quality. The core historical universe is sufficiently clean for the current analysis, but broader or more liquid modern universes may contain extreme adjusted-return events that require additional validation. These events may reflect corporate actions, ticker transitions, units or adjustment-factor issues. We would flag, inspect and either correct or handle them through robustness checks before using broader universes for final network claims.

The fourth limitation concerns sector classification. The macro-sector and subsector map is an initial classification. It is adequate for testing whether within-sector correlations exceed between-sector correlations, but it should be reviewed ticker by ticker before submission. This is especially important for holding companies, units, firms with multiple share classes and companies with mixed exposure across sectors.

Finally, the forecasting analysis is limited to three target assets, static network features and simple neural architectures. Feature importance in tree models is useful but does not establish causal mechanisms. PMFG features show nonzero but modest importance, and the magnitude is small compared with lagged-volatility predictors. Future work should include rolling network features, a broader target asset set, richer temporal neural architectures and graph neural networks.

13. Conclusion

We examined the dependency structure of Brazilian equities traded on B3 from 1998 to 2025 using an econophysics workflow. The main findings can be summarized along four dimensions.

First, the stylized-fact and correlation analysis shows that Brazilian equities display signatures commonly associated with complex financial systems: heavy-tailed return behaviour, volatility clustering, persistent absolute-return dependence and a moderate but positive average cross-correlation ($\bar{\rho} \approx 0.24$). Within-sector correlations are larger than between-sector correlations on average, confirming that sectoral organization is reflected in the dependency matrix.

Second, Random Matrix Theory decomposes the correlation matrix. Five eigenvalues exceed the Marčenko–Pastur upper bound, with the largest eigenvalue accounting for approximately 37.3% of the trace of the correlation matrix. The leading component captures a market-mode component, while the remaining eigenvectors outside the random band suggest commodity, financial and utility-related group structures. The reconstruction checks indicate numerical

accuracy within floating-point precision, and the group-mode matrix retains localized dependency patterns after the dominant market component is removed.

Third, financial network analysis through MST and PMFG representations shows that RMT filtering changes the structure of asset dependencies. The group-mode PMFG has lower mean edge correlation but higher clustering than the original PMFG, indicating stronger local clustering after market-mode removal. Hub rankings shift from networks dominated by broad market co-movement to networks where different assets occupy central positions after excluding the leading component. The PMFG retains 168 edges with triangles and 4-cliques, leaving local graph structures that are absent from the MST and useful for subsector-level analysis.

Fourth, the volatility forecasting experiment gives preliminary evidence that network topology can complement standard volatility predictors. Random Forest with network features obtains the best average QLIKE among the validation-selected models at the 5-day horizon, while Ridge regression with network features is strongest by QLIKE at the 20-day horizon. The CNN-1D model is competitive in MAE, RMSE and out-of-sample R^2 , but the neural extension does not uniformly dominate across metrics. Feature-importance analysis shows that lagged volatility features dominate, while selected PMFG centralities contribute smaller but nonzero predictive information.

We view the net contribution of network features as incremental: they do not replace lagged-volatility dynamics, but they add structural information from the cross-sectional geometry of the market. Future work should extend this framework by constructing rolling RMT and network features for real-time forecasting, expanding the target asset set and evaluating richer temporal neural network and graph neural network architectures.

A. Supplementary figures

Additional figures and tables generated by the computational workflow are retained as supplementary material. These include asset-universe tables, return summaries, correlation matrices, RMT eigenvector loadings, network edge lists, clique lists, centrality rankings and detailed forecasting results by asset, horizon, model and feature set.

B. Supplementary tables

Additional figures and tables generated by the computational workflow are retained as supplementary material. These include asset-universe tables, return summaries, correlation matrices, RMT eigenvector loadings, network edge lists, clique lists, centrality rankings and detailed forecasting results by asset, horizon, model and feature set.

Funding

This research did not receive any specific grant from funding agencies in the public, commercial, or not-for-profit sectors.

Declaration of competing interest

The authors declare that they have no known competing financial interests or personal relationships that could have appeared to influence the work reported in this paper.

Data availability

The workflow uses locally processed B3/BovDBv2 daily adjusted price records. The public BovDBv2 dataset is available through Harvard Dataverse under DOI 10.7910/DVN/TMB4IG. Derived scripts and reproducibility materials will be made available in a public repository upon publication, subject to redistribution constraints for raw market data.

Declaration of generative AI and AI-assisted technologies in the manuscript preparation process

During the preparation of this work, the authors used ChatGPT for language refinement, structural organization and code-review support. After using this tool, the authors reviewed and edited the content as needed and take full responsibility for the content of the published article.

CRedit authorship contribution statement

Andrey V. S. Souza: Conceptualization, Methodology, Software, Formal analysis, Writing - original draft. **Bruno L. Dalmazo:** Methodology, Supervision, Writing - review and editing. **Giancarlo Lucca:** Methodology, Validation, Writing - review and editing. **Eduardo N. Borges:** Supervision, Methodology, Writing - review and editing. **Richard F. Pinto:** Software, Data curation, Validation. **Fabian C. Cardoso:** Data curation, Resources, Writing - review and editing. **Viviane L. D. de Mattos:** Formal analysis, Supervision, Writing - review and editing. **Rafael A. Berri:** Supervision, Project administration, Writing - review and editing.

References

- [1] R. N. Mantegna, H. E. Stanley, *Introduction to Econophysics: Correlations and Complexity in Finance*, Cambridge University Press, Cambridge, 1999.
- [2] R. Cont, Empirical properties of asset returns: stylized facts and statistical issues, *Quantitative Finance* 1 (2) (2001) 223–236. doi:10.1080/713665670.
- [3] M. Raddant, T. Di Matteo, A look at financial dependencies by means of econophysics and financial economics, *Journal of Economic Interaction and Coordination* 18 (4) (2023) 701–734. doi:10.1007/s11403-023-00395-y.
- [4] M. Billio, M. Getmansky, A. W. Lo, L. Pelizzon, Econometric measures of connectedness and systemic risk in the finance and insurance sectors, *Journal of Financial Economics* 104 (3) (2012) 535–559. doi:10.1016/j.jfineco.2011.12.010.
- [5] L. S. Junior, I. T. P. Franca, Correlation of financial markets in times of crisis, *Physica A: Statistical Mechanics and its Applications* 391 (1–2) (2012) 187–208. doi:10.1016/j.physa.2011.07.023.
- [6] Z. Zheng, B. Podobnik, L. Feng, B. Li, Changes in cross-correlations as an indicator for systemic risk, *Scientific Reports* 2 (2012) 888. doi:10.1038/srep00888.
- [7] B. M. Tabak, D. O. Cajueiro, M. V. Serra, Topology of the Brazilian interbank network, *Physica A: Statistical Mechanics and its Applications* 389 (22) (2010) 5277–5283. doi:10.1016/j.physa.2010.07.037.
- [8] T. C. Silva, S. R. S. de Souza, B. M. Tabak, Structure and dynamics of the global financial network, *Chaos, Solitons & Fractals* 88 (2016) 218–234. doi:10.1016/j.chaos.2016.01.023.
- [9] T. Stosić, L. Telesca, T. A. E. Ferreira, B. D. Stosic, Correlations and cross-correlations in the Brazilian agrarian commodities and stocks, *Physica A: Statistical Mechanics and its Applications* 437 (2015) 248–262. doi:10.1016/j.physa.2015.05.091.
- [10] E. J. A. L. Pereira, M. F. da Silva, H. B. B. Pereira, Multiscale network for twenty stock markets using DCCA, *Physica A: Statistical Mechanics and its Applications* 529 (2019) 121542. doi:10.1016/j.physa.2019.121542.
- [11] M. A. Leonel Caetano, T. Yoneyama, Complex network analysis of the Brazilian stock market, *Expert Systems with Applications* 196 (2022) 116680. doi:10.1016/j.eswa.2022.116680.
- [12] V. A. Marčenko, L. A. Pastur, Distribution of eigenvalues for some sets of random matrices, *Mathematics of the USSR-Sbornik* 1 (4) (1967) 457–483. doi:10.1070/SM1967v001n04ABEH001994.
- [13] L. Laloux, P. Cizeau, J.-P. Bouchaud, M. Potters, Noise dressing of financial correlation matrices, *Physical Review Letters* 83 (7) (1999) 1467–1470. doi:10.1103/PhysRevLett.83.1467.
- [14] V. Plerou, P. Gopikrishnan, B. Rosenow, L. A. N. Amaral, T. Guhr, H. E. Stanley, Random matrix approach to cross correlations in financial data, *Physical Review E* 65 (6) (2002) 066126. doi:10.1103/PhysRevE.65.066126.
- [15] J. Bun, J.-P. Bouchaud, M. Potters, Cleaning large correlation matrices: Tools from random matrix theory, *Physics Reports* 666 (2017) 1–109. doi:10.1016/j.physrep.2016.10.005.
- [16] R. N. Mantegna, Hierarchical structure in financial markets, *The European Physical Journal B* 11 (1) (1999) 193–197. doi:10.1007/s100510050929.
- [17] M. Tumminello, T. Aste, T. Di Matteo, R. N. Mantegna, A tool for filtering information in complex systems, *Proceedings of the National Academy of Sciences* 102 (30) (2005) 10421–10426. doi:10.1073/pnas.0500298102.
- [18] M. Tumminello, F. Lillo, R. N. Mantegna, Correlation, hierarchies, and networks in financial markets, *Journal of Economic Behavior & Organization* 75 (1) (2010) 40–58. doi:10.1016/j.jebo.2010.01.004.
- [19] F. Pozzi, T. Di Matteo, T. Aste, Spread of risk across financial markets: better to invest in the peripheries, *Scientific Reports* 3 (2013) 1665. doi:10.1038/srep01665.
- [20] T. Aste, W. Shaw, T. Di Matteo, Correlation structure and dynamics in volatile markets, *New Journal of Physics* 12 (8) (2010) 085009. doi:10.1088/1367-2630/12/8/085009.
- [21] L. Laloux, P. Cizeau, M. Potters, J.-P. Bouchaud, Random matrix theory and financial correlations, *International Journal of Theoretical and Applied Finance* 3 (3) (2000) 391–397. doi:10.1142/S0219024900000255.
- [22] T. Guhr, A. Müller-Groeling, H. A. Weidenmüller, Random-matrix theories in quantum physics: common concepts, *Physics Reports* 299 (4–6) (1998) 189–425. doi:10.1016/S0370-1573(97)00088-4.
- [23] J.-P. Onnela, A. Chakraborti, K. Kaski, J. Kertész, A. Kanto, Dynamics of market correlations: Taxonomy and portfolio analysis, *Physical Review E* 68 (5) (2003) 056110. doi:10.1103/PhysRevE.68.056110.
- [24] G. Bonanno, G. Caldarelli, F. Lillo, S. Miccichè, N. Vandewalle, R. N. Mantegna, Networks of equities in financial markets, *The European Physical Journal B* 38 (2) (2004) 363–371. doi:10.1140/epjb/e2004-00129-6.
- [25] D.-M. Song, M. Tumminello, W.-X. Zhou, R. N. Mantegna, Evolution of worldwide stock markets, correlation structure, and correlation-based graphs, *Physical Review E* 84 (2) (2011) 026108. doi:10.1103/PhysRevE.84.026108.

- [26] G. P. Massara, T. Di Matteo, T. Aste, Network filtering for big data: Triangulated maximally filtered graph, *Journal of Complex Networks* 5 (2) (2017) 161–178. doi:10.1093/comnet/cnw015.
- [27] D. Y. Kenett, M. Tumminello, A. Madi, G. Gur-Gershgoren, R. N. Mantegna, E. Ben-Jacob, Dominating clasp of the financial sector revealed by partial correlation analysis of the stock market, *PLoS ONE* 5 (12) (2010) e15032. doi:10.1371/journal.pone.0015032.
- [28] G.-J. Wang, C. Xie, H. E. Stanley, Correlation structure and evolution of world stock markets: Evidence from Pearson and partial correlation-based networks, *Computational Economics* 51 (3) (2018) 607–635. doi:10.1007/s10614-016-9627-7.
- [29] T. G. Andersen, T. Bollerslev, F. X. Diebold, P. Labys, Modeling and forecasting realized volatility, *Econometrica* 71 (2) (2003) 579–625. doi:10.1111/1468-0262.00418.
- [30] S.-H. Poon, C. W. J. Granger, Forecasting volatility in financial markets: A review, *Journal of Economic Literature* 41 (2) (2003) 478–539. doi:10.1257/jel.41.2.478.
- [31] A. J. Patton, Volatility forecast comparison using imperfect volatility proxies, *Journal of Econometrics* 160 (1) (2011) 246–256. doi:10.1016/j.jeconom.2010.03.034.
- [32] L. Breiman, Random forests, *Machine Learning* 45 (1) (2001) 5–32. doi:10.1023/A:1010933404324.
- [33] J. H. Friedman, Greedy function approximation: A gradient boosting machine, *Annals of Statistics* 29 (5) (2001) 1189–1232. doi:10.1214/aos/1013203451.
- [34] A. E. Hoerl, R. W. Kennard, Ridge regression: Biased estimation for nonorthogonal problems, *Technometrics* 12 (1) (1970) 55–67. doi:10.1080/00401706.1970.10488634.
- [35] A. Bucci, Realized volatility forecasting with neural networks, *Journal of Financial Econometrics* 18 (3) (2020) 502–531. doi:10.1093/jjfinec/nbaa008.
- [36] K. Christensen, M. Siggaard, B. Veliyev, A machine learning approach to volatility forecasting, *Journal of Financial Econometrics* 21 (5) (2023) 1680–1727. doi:10.1093/jjfinec/nbac020.
- [37] D. Zhang, D. C. Broadstock, Financial network topology and market risk: A machine learning perspective, *International Review of Financial Analysis* 87 (2023) 102604. doi:10.1016/j.irfa.2023.102604.
- [38] F. C. Cardoso, J. A. V. Malska, P. J. Ramiro, G. Lucca, E. N. Borges, V. L. D. d. Mattos, R. A. Berri, BovDB: a data set of stock prices of all companies in B3 from 1995 to 2020, *Journal of Information and Data Management* 13 (1). doi:10.5753/jidm.2022.2345.
- [39] A. S. Souza, G. Lucca, E. N. Borges, F. C. Cardoso, B. L. Dalmazo, R. Berri, Dataset for Intraday Analysis of B3 stock prices, *Harvard Dataverse*, v1 (2024). doi:10.7910/DVN/TMB4IG.
URL <https://doi.org/10.7910/DVN/TMB4IG>
- [40] R. R. Sokal, F. J. Rohlf, The comparison of dendrograms by objective methods, *Taxon* 11 (2) (1962) 33–40. doi:10.2307/1217208.